

# autogluon-2

November 17, 2024

- 1 Using AWS Autogluon to predict rent prices in Canada based on city, province, latitude, longitude, lease\_term, type, price, beds, baths, sq\_feet, furnishing, availability\_date, smoking, and whether cats and dogs are allowed

```
[1]: # !pip install autogluon
```

```
[3]: import pandas as pd
import numpy as np
```

```
[4]: df = pd.read_csv("rentfaster.csv")
```

```
[5]: df.head()
```

```
[5]:
```

	rentfaster_id	city	province	address	latitude	longitude	\
0	468622	Airdrie	Alberta	69 Gateway Dr NE	51.305962	-114.012515	
1	468622	Airdrie	Alberta	69 Gateway Dr NE	51.305962	-114.012515	
2	468622	Airdrie	Alberta	69 Gateway Dr NE	51.305962	-114.012515	
3	468622	Airdrie	Alberta	69 Gateway Dr NE	51.305962	-114.012515	
4	468622	Airdrie	Alberta	69 Gateway Dr NE	51.305962	-114.012515	

	lease_term	type	price	beds	baths	sq_feet	\
0	Long Term	Townhouse	2495.0	2 Beds	2.5	1403	
1	Long Term	Townhouse	2695.0	3 Beds	2.5	1496	
2	Long Term	Townhouse	2295.0	2 Beds	2.5	1180	
3	Long Term	Townhouse	2095.0	2 Beds	2.5	1403	
4	Long Term	Townhouse	2495.0	2 Beds	2.5	1403	

	link	furnishing	\
0	/ab/airdrie/rentals/townhouse/2-bedrooms/pet-f...	Unfurnished	
1	/ab/airdrie/rentals/townhouse/2-bedrooms/pet-f...	Unfurnished	
2	/ab/airdrie/rentals/townhouse/2-bedrooms/pet-f...	Unfurnished	
3	/ab/airdrie/rentals/townhouse/2-bedrooms/pet-f...	Unfurnished	
4	/ab/airdrie/rentals/townhouse/2-bedrooms/pet-f...	Unfurnished	

	availability_date	smoking	cats	dogs
--	-------------------	---------	------	------

```

0      Immediate Non-Smoking True True
1      Immediate Non-Smoking True True
2      Immediate Non-Smoking True True
3      November 18 Non-Smoking True True
4      Immediate Non-Smoking True True

```

```
[6]: df.columns
```

```
[6]: Index(['rentfaster_id', 'city', 'province', 'address', 'latitude', 'longitude',
          'lease_term', 'type', 'price', 'beds', 'baths', 'sq_feet', 'link',
          'furnishing', 'availability_date', 'smoking', 'cats', 'dogs'],
          dtype='object')
```

```
[7]: # Remove unnecessary columns
df = df.drop(axis=1, columns = ["address", "link", "rentfaster_id"])
```

```
[8]: df.head()
```

```
[8]:
```

	city	province	latitude	longitude	lease_term	type	price	\
0	Airdrie	Alberta	51.305962	-114.012515	Long Term	Townhouse	2495.0	
1	Airdrie	Alberta	51.305962	-114.012515	Long Term	Townhouse	2695.0	
2	Airdrie	Alberta	51.305962	-114.012515	Long Term	Townhouse	2295.0	
3	Airdrie	Alberta	51.305962	-114.012515	Long Term	Townhouse	2095.0	
4	Airdrie	Alberta	51.305962	-114.012515	Long Term	Townhouse	2495.0	

	beds	baths	sq_feet	furnishing	availability_date	smoking	cats	\
0	2 Beds	2.5	1403	Unfurnished	Immediate	Non-Smoking	True	
1	3 Beds	2.5	1496	Unfurnished	Immediate	Non-Smoking	True	
2	2 Beds	2.5	1180	Unfurnished	Immediate	Non-Smoking	True	
3	2 Beds	2.5	1403	Unfurnished	November 18	Non-Smoking	True	
4	2 Beds	2.5	1403	Unfurnished	Immediate	Non-Smoking	True	

	dogs
0	True
1	True
2	True
3	True
4	True

```
[9]: from autogluon.tabular import TabularDataset, TabularPredictor
```

```
[10]: # Regressing for the price
target = "price"
```

```
[11]: train_data = TabularDataset(df)
```

```
[12]: # Sample 70% randomly for the train data
      subsample_size = int(0.7*len(df))
```

```
[13]: train_data = train_data.sample(n=subsample_size, random_state=0)
      train_data.head()
```

```
[13]:
```

	city	province	latitude	longitude	lease_term	type	\
12512	Calgary	Alberta	51.032613	-114.062190	Long Term	Condo Unit	
24216	Montréal	Quebec	45.505681	-73.563915	Long Term	Apartment	
13161	Calgary	Alberta	50.859881	-114.078010	Long Term	Basement	
2415	Calgary	Alberta	51.134793	-113.949708	Long Term	Condo Unit	
7519	Edmonton	Alberta	53.544671	-113.577309	12 months	House	

	price	beds	baths	sq_feet	furnishing	availability_date	\
12512	2150.0	1 Bed	1	763	Unfurnished	July 01	
24216	3390.0	2 Beds	1	1058	Unfurnished	Immediate	
13161	1300.0	1 Bed	1	700	Unfurnished	July 01	
2415	2150.0	2 Beds	2	980	Unfurnished	Immediate	
7519	1700.0	2 Beds	2	800	Unfurnished	July 01	

	smoking	cats	dogs
12512	Non-Smoking	False	False
24216	Non-Smoking	True	True
13161	Non-Smoking	True	True
2415	Non-Smoking	False	False
7519	Non-Smoking	True	True

```
[13]: # Training a fast model
      predictor_price = TabularPredictor(label=target).fit(train_data, time_limit=60)
```

No path specified. Models will be saved in: "AutogluonModels/ag-20241117\_161814"  
Verbosity: 2 (Standard Logging)

===== System Info =====

```
AutoGluon Version: 1.1.1
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP PREEMPT_DYNAMIC Thu Jun 27 21:05:47 UTC 2024
CPU Count: 2
Memory Avail: 10.88 GB / 12.67 GB (85.9%)
Disk Space Avail: 65.91 GB / 107.72 GB (61.2%)
=====
```

No presets specified! To achieve strong results with AutoGluon, it is recommended to use the available presets.

Recommended Presets (For more details refer to  
<https://auto.gluon.ai/stable/tutorials/tabular/tabular-essentials.html#presets>):  
 presets='best\_quality' : Maximize accuracy. Default time\_limit=3600.

```

    presets='high_quality' : Strong accuracy with fast inference speed.
Default time_limit=3600.
    presets='good_quality' : Good accuracy with very fast inference speed.
Default time_limit=3600.
    presets='medium_quality' : Fast training time, ideal for initial
prototyping.
Beginning AutoGluon training ... Time limit = 60s
AutoGluon will save models to "AutogluonModels/ag-20241117_161814"
Train Data Rows:      18039
Train Data Columns: 14
Label Column:        price
AutoGluon infers your prediction problem is: 'regression' (because dtype of
label-column == float and many unique label-values observed).
    Label info (max, min, mean, stddev): (29990.0, 0.0, 2144.43951,
972.07588)
    If 'regression' is not the correct problem_type, please manually specify
the problem_type parameter during Predictor init (You may specify problem_type
as one of: ['binary', 'multiclass', 'regression', 'quantile'])
Problem Type:        regression
Preprocessing data ...
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
    Available Memory:                11156.12 MB
    Train Data (Original) Memory Usage: 12.43 MB (0.1% of available memory)
    Inferring data type of each feature based on column values. Set
feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
        Fitting CategoryFeatureGenerator...
        Fitting CategoryMemoryMinimizeFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
Types of features in original data (raw dtype, special dtypes):
    ('float', []) : 2 | ['latitude', 'longitude']
    ('object', []) : 12 | ['city', 'province', 'lease_term', 'type',
'beds', ...]
Types of features in processed data (raw dtype, special dtypes):
    ('category', []) : 12 | ['city', 'province', 'lease_term',
'type', 'beds', ...]
    ('float', []) : 2 | ['latitude', 'longitude']
0.6s = Fit runtime
14 features in original data used to generate 14 features in processed

```

data.

Train Data (Processed) Memory Usage: 0.52 MB (0.0% of available memory)

Data preprocessing and feature engineering runtime = 0.63s ...

AutoGluon will gauge predictive performance using evaluation metric:

'root\_mean\_squared\_error'

This metric's sign has been flipped to adhere to being higher\_is\_better.

The metric score can be multiplied by -1 to get the metric value.

To change this, specify the eval\_metric parameter of Predictor()

Automatically generating train/validation split with holdout\_frac=0.1, Train

Rows: 16235, Val Rows: 1804

User-specified model hyperparameters to be fit:

```
{
    'NN_TORCH': {},
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
    'GBMLarge': {},
    'CAT': {},
    'XGB': {},
    'FASTAI': {},
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
        'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
        {'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
        {'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
        'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
        'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
        {'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
        {'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
        'problem_types': ['regression', 'quantile']}}],
    'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
        {'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
```

Fitting 11 L1 models ...

Fitting model: KNeighborsUnif ... Training model for up to 59.37s of the 59.37s of remaining time.

-683.1182 = Validation score (-root\_mean\_squared\_error)

2.96s = Training runtime

0.02s = Validation runtime

Fitting model: KNeighborsDist ... Training model for up to 56.38s of the 56.37s of remaining time.

-612.811 = Validation score (-root\_mean\_squared\_error)

0.02s = Training runtime

0.02s = Validation runtime

Fitting model: LightGBMXT ... Training model for up to 56.33s of the 56.33s of remaining time.

/usr/local/lib/python3.10/dist-packages/dask/dataframe/\_\_init\_\_.py:42:

FutureWarning:

Dask dataframe query planning is disabled because dask-expr is not installed.

You can install it with `pip install dask[dataframe]` or `conda install dask`. This will raise in a future version.

```
warnings.warn(msg, FutureWarning)

[1000] valid_set's rmse: 423.81
[2000] valid_set's rmse: 415.834
[3000] valid_set's rmse: 410.2
[4000] valid_set's rmse: 407.48
[5000] valid_set's rmse: 404.523
[6000] valid_set's rmse: 402.728
[7000] valid_set's rmse: 400.988
[8000] valid_set's rmse: 399.787
[9000] valid_set's rmse: 399.434
[10000] valid_set's rmse: 398.293

-398.293          = Validation score    (-root_mean_squared_error)
30.04s           = Training    runtime
10.43s           = Validation runtime
Fitting model: LightGBM ... Training model for up to 14.44s of the 14.44s of
remaining time.

[1000] valid_set's rmse: 378.305
[2000] valid_set's rmse: 365.213
[3000] valid_set's rmse: 360.246
[4000] valid_set's rmse: 357.389
[5000] valid_set's rmse: 355.923

Ran out of time, early stopping on iteration 5875. Best iteration is:
[5863] valid_set's rmse: 354.889
-354.8894         = Validation score    (-root_mean_squared_error)
15.51s           = Training    runtime
1.97s            = Validation runtime
Fitting model: WeightedEnsemble_L2 ... Training model for up to 59.37s of the
-4.59s of remaining time.
Ensemble Weights: {'LightGBM': 0.737, 'KNeighborsDist': 0.158,
'LightGBMXT': 0.105}
-341.2729         = Validation score    (-root_mean_squared_error)
0.02s            = Training    runtime
0.0s             = Validation runtime
AutoGluon training complete, total runtime = 64.66s ... Best model:
WeightedEnsemble_L2 | Estimated inference throughput: 145.3 rows/s (1804 batch
size)
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("AutogluonModels/ag-20241117_161814")

[14]: # Creating test data for the rows that have not been sampled by the train data
test_data = TabularDataset(df.drop(train_data.index))
y_test = test_data[target]
```

```
[15]: # predictor = TabularPredictor.load("agModels-predictprice")

# Predicting rent prices for the test data
y_pred = predictor_price.predict(test_data)
print("Predictions: \n", y_pred)
perf = predictor_price.evaluate_predictions(y_true=y_test, y_pred=y_pred,
    ↪ auxiliary_metrics=True)
```

Predictions:

```
10      2736.225098
11      2212.453857
13      2559.374023
19      2113.462646
24      2080.962402
```

...

```
25748    1267.991943
25755    1058.244019
25762     945.008179
25763    1113.939209
25764     945.008179
```

Name: price, Length: 7732, dtype: float32

```
[16]: # Looking at the performance
perf
```

```
[16]: {'root_mean_squared_error': -394.74039693281304,
      'mean_squared_error': -155819.9809706748,
      'mean_absolute_error': -183.9261484073692,
      'r2': 0.8358683478621021,
      'pearsonr': 0.9144924898179707,
      'median_absolute_error': -98.54681396484375}
```

```
[18]: # retraining the model using best_quality and time limit of 10 minutes,
    ↪ focusing on MAE
predictor_price = TabularPredictor(label=target, path="agModels-predictprice",
    eval_metric="root_mean_squared_error").
    ↪ fit(train_data,
        presets="best_quality", time_limit=600)
```

Verbosity: 2 (Standard Logging)

===== System Info =====

```
AutoGluon Version: 1.1.1
Python Version: 3.10.12
Operating System: Linux
Platform Machine: x86_64
Platform Version: #1 SMP PREEMPT_DYNAMIC Thu Jun 27 21:05:47 UTC 2024
CPU Count: 2
```

Memory Avail: 10.71 GB / 12.67 GB (84.5%)  
Disk Space Avail: 65.87 GB / 107.72 GB (61.1%)

=====

Presets specified: ['best\_quality']

Setting dynamic\_stacking from 'auto' to True. Reason: Enable dynamic\_stacking when use\_bag\_holdout is disabled. (use\_bag\_holdout=False)

Stack configuration (auto\_stack=True): num\_stack\_levels=1, num\_bag\_folds=8, num\_bag\_sets=1

DyStack is enabled (dynamic\_stacking=True). AutoGluon will try to determine whether the input data is affected by stacked overfitting and enable or disable stacking as a consequence.

This is used to identify the optimal `num\_stack\_levels` value. Copies of AutoGluon will be fit on subsets of the data. Then holdout validation data is used to detect stacked overfitting.

Running DyStack for up to 150s of the 600s of remaining time (25%).

Running DyStack sub-fit in a ray process to avoid memory leakage.

Enabling ray logging (enable\_ray\_logging=True). Specify

`ds\_args={'enable\_ray\_logging': False}` if you experience logging issues.

2024-11-17 16:21:56,940 INFO worker.py:1743 -- Started a local Ray instance.

View the dashboard at [127.0.0.1:8265](http://127.0.0.1:8265)

Context path: "agModels-predictprice/ds\_sub\_fit/sub\_fit\_ho"

(\_dystack pid=4697) Running DyStack sub-fit ...

(\_dystack pid=4697) Beginning AutoGluon training ... Time limit = 141s

(\_dystack pid=4697) AutoGluon will save models to "agModels-predictprice/ds\_sub\_fit/sub\_fit\_ho"

(\_dystack pid=4697) Train Data Rows: 16034

(\_dystack pid=4697) Train Data Columns: 14

(\_dystack pid=4697) Label Column: price

(\_dystack pid=4697) Problem Type: regression

(\_dystack pid=4697) Preprocessing data ...

(\_dystack pid=4697) Using Feature Generators to preprocess the data ...

(\_dystack pid=4697) Fitting AutoMLPipelineFeatureGenerator...

(\_dystack pid=4697) Available Memory: 10565.92 MB

(\_dystack pid=4697) Train Data (Original) Memory Usage: 11.04 MB

(0.1% of available memory)

(\_dystack pid=4697) Inferring data type of each feature based on column values. Set feature\_metadata\_in to manually specify special dtypes of the features.

(\_dystack pid=4697) Stage 1 Generators:

(\_dystack pid=4697) Fitting AsTypeFeatureGenerator...

(\_dystack pid=4697) Stage 2 Generators:

(\_dystack pid=4697) Fitting FillNaFeatureGenerator...

(\_dystack pid=4697) Stage 3 Generators:

(\_dystack pid=4697) Fitting IdentityFeatureGenerator...

(\_dystack pid=4697) Fitting CategoryFeatureGenerator...

(\_dystack pid=4697) Fitting

CategoryMemoryMinimizeFeatureGenerator...

(\_dystack pid=4697) Stage 4 Generators:



```

(_dystack pid=4697)          Fitting DropUniqueFeatureGenerator...
(_dystack pid=4697)    Stage 5 Generators:
(_dystack pid=4697)          Fitting
DropDuplicatesFeatureGenerator...
(_dystack pid=4697)    Types of features in original data (raw dtype,
special dtypes):
(_dystack pid=4697)          ('float', []) : 2 | ['latitude',
'longitude']
(_dystack pid=4697)          ('object', []) : 12 | ['city',
'province', 'lease_term', 'type', 'beds', ...]
(_dystack pid=4697)    Types of features in processed data (raw dtype,
special dtypes):
(_dystack pid=4697)          ('category', []) : 12 | ['city',
'province', 'lease_term', 'type', 'beds', ...]
(_dystack pid=4697)          ('float', []) : 2 | ['latitude',
'longitude']
(_dystack pid=4697)    0.3s = Fit runtime
(_dystack pid=4697)    14 features in original data used to generate 14
features in processed data.
(_dystack pid=4697)    Train Data (Processed) Memory Usage: 0.46 MB
(0.0% of available memory)
(_dystack pid=4697) Data preprocessing and feature engineering runtime
= 0.28s ...
(_dystack pid=4697) AutoGluon will gauge predictive performance using
evaluation metric: 'root_mean_squared_error'
(_dystack pid=4697)    This metric's sign has been flipped to adhere to
being higher_is_better. The metric score can be multiplied by -1 to get the
metric value.
(_dystack pid=4697)    To change this, specify the eval_metric
parameter of Predictor()
(_dystack pid=4697) Large model count detected (112 configs) ... Only
displaying the first 3 models of each family. To see all, set `verbosity=3`.
(_dystack pid=4697) User-specified model hyperparameters to be fit:
(_dystack pid=4697) {
(_dystack pid=4697)     'NN_TORCH': [{}, {'activation': 'elu',
'dropout_prob': 0.10077639529843717, 'hidden_size': 108, 'learning_rate':
0.002735937344002146, 'num_layers': 4, 'use_batchnorm': True, 'weight_decay':
1.356433327634438e-12, 'ag_args': {'name_suffix': '_r79', 'priority': -2}},
{'activation': 'elu', 'dropout_prob': 0.11897478034205347, 'hidden_size': 213,
'learning_rate': 0.0010474382260641949, 'num_layers': 4, 'use_batchnorm': False,
'weight_decay': 5.594471067786272e-10, 'ag_args': {'name_suffix': '_r22',
'priority': -7}}],
(_dystack pid=4697)     'GBM': [{ 'extra_trees': True, 'ag_args':
{'name_suffix': 'XT'}}, {}, 'GBMLarge'],
(_dystack pid=4697)     'CAT': [{}, {'depth': 6, 'grow_policy':
'SymmetricTree', 'l2_leaf_reg': 2.1542798306067823, 'learning_rate':
0.06864209415792857, 'max_ctr_complexity': 4, 'one_hot_max_size': 10, 'ag_args':
{'name_suffix': '_r177', 'priority': -1}}, {'depth': 8, 'grow_policy':

```

```

'Depthwise', 'l2_leaf_reg': 2.7997999596449104, 'learning_rate':
0.031375015734637225, 'max_ctr_complexity': 2, 'one_hot_max_size': 3, 'ag_args':
{'name_suffix': '_r9', 'priority': -5}},
(_dystack pid=4697) 'XGB': [{}, {'colsample_bytree':
0.6917311125174739, 'enable_categorical': False, 'learning_rate':
0.018063876087523967, 'max_depth': 10, 'min_child_weight': 0.6028633586934382,
'ag_args': {'name_suffix': '_r33', 'priority': -8}}, {'colsample_bytree':
0.6628423832084077, 'enable_categorical': False, 'learning_rate':
0.08775715546881824, 'max_depth': 5, 'min_child_weight': 0.6294123374222513,
'ag_args': {'name_suffix': '_r89', 'priority': -16}}],
(_dystack pid=4697) 'FASTAI': [{}, {'bs': 256, 'emb_drop':
0.5411770367537934, 'epochs': 43, 'layers': [800, 400], 'lr':
0.01519848858318159, 'ps': 0.23782946566604385, 'ag_args': {'name_suffix':
'_r191', 'priority': -4}}, {'bs': 2048, 'emb_drop': 0.05070411322605811,
'epochs': 29, 'layers': [200, 100], 'lr': 0.08974235041576624, 'ps':
0.10393466140748028, 'ag_args': {'name_suffix': '_r102', 'priority': -11}}],
(_dystack pid=4697) 'RF': [{'criterion': 'gini', 'ag_args':
{'name_suffix': 'Gini', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'entropy', 'ag_args': {'name_suffix': 'Entr', 'problem_types':
['binary', 'multiclass']}}, {'criterion': 'squared_error', 'ag_args':
{'name_suffix': 'MSE', 'problem_types': ['regression', 'quantile']}}],
(_dystack pid=4697) 'XT': [{'criterion': 'gini', 'ag_args':
{'name_suffix': 'Gini', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'entropy', 'ag_args': {'name_suffix': 'Entr', 'problem_types':
['binary', 'multiclass']}}, {'criterion': 'squared_error', 'ag_args':
{'name_suffix': 'MSE', 'problem_types': ['regression', 'quantile']}}],
(_dystack pid=4697) 'KNN': [{'weights': 'uniform', 'ag_args':
{'name_suffix': 'Unif'}}], {'weights': 'distance', 'ag_args': {'name_suffix':
'Dist'}}],
(_dystack pid=4697) }
(_dystack pid=4697) AutoGluon will fit 2 stack levels (L1 to L2) ...
(_dystack pid=4697) Fitting 108 L1 models ...
(_dystack pid=4697) Fitting model: KNeighborsUnif_BAG_L1 ... Training
model for up to 94.07s of the 141.13s of remaining time.
(_dystack pid=4697) -802.1015 = Validation score
(-root_mean_squared_error)
(_dystack pid=4697) 0.02s = Training runtime
(_dystack pid=4697) 0.06s = Validation runtime
(_dystack pid=4697) Fitting model: KNeighborsDist_BAG_L1 ... Training
model for up to 91.93s of the 138.98s of remaining time.
(_dystack pid=4697) -685.5784 = Validation score
(-root_mean_squared_error)
(_dystack pid=4697) 0.03s = Training runtime
(_dystack pid=4697) 0.05s = Validation runtime
(_dystack pid=4697) Fitting model: LightGBMXT_BAG_L1 ... Training model
for up to 91.83s of the 138.88s of remaining time.
(_dystack pid=4697) Fitting 8 child models (S1F1 - S1F8) | Fitting
with ParallelLocalFoldFittingStrategy (2 workers, per: cpus=1, gpus=0,

```

```

memory=0.06%)
(_ray_fit pid=4861) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning:
(_ray_fit pid=4861) Dask dataframe query planning is disabled because
dask-expr is not installed.
(_ray_fit pid=4861)
(_ray_fit pid=4861) You can install it with `pip install
dask[dataframe]` or `conda install dask`.
(_ray_fit pid=4861) This will raise in a future version.
(_ray_fit pid=4861)
(_ray_fit pid=4861) warnings.warn(msg, FutureWarning)

(_ray_fit pid=4861) [1000]      valid_set's rmse: 475.131
(_ray_fit pid=4861) [3000]      valid_set's rmse: 452.806 [repeated
4x across cluster] (Ray deduplicates logs by default. Set RAY_DEDUP_LOGS=0 to
disable log deduplication, or see https://docs.ray.io/en/master/ray-
observability/ray-logging.html#log-deduplication for more options.)

(_ray_fit pid=5050) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning: [repeated 2x across
cluster]
(_ray_fit pid=5050) Dask dataframe query planning is disabled because
dask-expr is not installed. [repeated 2x across cluster]
(_ray_fit pid=5050) [repeated 4x across cluster]
(_ray_fit pid=5050) You can install it with `pip install
dask[dataframe]` or `conda install dask`. [repeated 2x across cluster]
(_ray_fit pid=5050) This will raise in a future version. [repeated
2x across cluster]
(_ray_fit pid=5050) warnings.warn(msg, FutureWarning) [repeated
2x across cluster]

(_ray_fit pid=5050) [1000]      valid_set's rmse: 454.321 [repeated
2x across cluster]
(_ray_fit pid=5050) [3000]      valid_set's rmse: 451.699 [repeated
4x across cluster]

(_ray_fit pid=5216) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning: [repeated 2x across
cluster]
(_ray_fit pid=5216) Dask dataframe query planning is disabled because
dask-expr is not installed. [repeated 2x across cluster]
(_ray_fit pid=5216) [repeated 4x across cluster]
(_ray_fit pid=5216) You can install it with `pip install
dask[dataframe]` or `conda install dask`. [repeated 2x across cluster]

```

```

(_ray_fit pid=5216) This will raise in a future version. [repeated
2x across cluster]
(_ray_fit pid=5216) warnings.warn(msg, FutureWarning) [repeated
2x across cluster]

(_ray_fit pid=5216) [1000] valid_set's rmse: 442.112 [repeated
2x across cluster]
(_ray_fit pid=5216) [2000] valid_set's rmse: 412.533 [repeated
2x across cluster]
(_ray_fit pid=5216) [4000] valid_set's rmse: 387.657 [repeated
4x across cluster]

(_ray_fit pid=5404) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning: [repeated 2x across
cluster]
(_ray_fit pid=5404) Dask dataframe query planning is disabled because
dask-expr is not installed. [repeated 2x across cluster]
(_ray_fit pid=5404) [repeated 4x across cluster]
(_ray_fit pid=5404) You can install it with `pip install
dask[dataframe]` or `conda install dask`. [repeated 2x across cluster]
(_ray_fit pid=5404) This will raise in a future version. [repeated
2x across cluster]
(_ray_fit pid=5404) warnings.warn(msg, FutureWarning) [repeated
2x across cluster]

(_ray_fit pid=5433) [1000] valid_set's rmse: 551.309 [repeated
2x across cluster]
(_ray_fit pid=5433) [3000] valid_set's rmse: 488.475 [repeated
2x across cluster]
(_ray_fit pid=5433) [6000] valid_set's rmse: 459.421 [repeated
3x across cluster]

(_dystack pid=4697) -484.9849 = Validation score
(-root_mean_squared_error)
(_dystack pid=4697) 117.28s = Training runtime
(_dystack pid=4697) 25.87s = Validation runtime
(_dystack pid=4697) Fitting model: WeightedEnsemble_L2 ... Training
model for up to 141.14s of the 12.97s of remaining time.
(_ray_fit pid=5433) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning:
(_ray_fit pid=5433) Dask dataframe query planning is disabled because
dask-expr is not installed.
(_ray_fit pid=5433) [repeated 2x across cluster]

```

```

(_ray_fit pid=5433) You can install it with `pip install
dask[dataframe]` or `conda install dask`.
(_ray_fit pid=5433) This will raise in a future version.
(_ray_fit pid=5433) warnings.warn(msg, FutureWarning)
(_dystack pid=4697) Ensemble Weights: {'LightGBMXT_BAG_L1': 0.769,
'KNeighborsDist_BAG_L1': 0.231}
(_dystack pid=4697) -460.7118 = Validation score
(-root_mean_squared_error)
(_dystack pid=4697) 0.02s = Training runtime
(_dystack pid=4697) 0.0s = Validation runtime
(_dystack pid=4697) Fitting 106 L2 models ...
(_dystack pid=4697) Fitting model: LightGBMXT_BAG_L2 ... Training model
for up to 12.93s of the 12.91s of remaining time.
(_dystack pid=4697) Fitting 8 child models (S1F1 - S1F8) | Fitting
with ParallelLocalFoldFittingStrategy (2 workers, per: cpus=1, gpus=0,
memory=0.08%)
(_ray_fit pid=5585) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning:
(_ray_fit pid=5585) Dask dataframe query planning is disabled because
dask-expr is not installed.
(_ray_fit pid=5585) You can install it with `pip install
dask[dataframe]` or `conda install dask`.
(_ray_fit pid=5585) This will raise in a future version.
(_ray_fit pid=5585) warnings.warn(msg, FutureWarning)
(_ray_fit pid=5585) [repeated 2x across cluster]
(_ray_fit pid=5693) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning: [repeated 2x across
cluster]
(_ray_fit pid=5693) Dask dataframe query planning is disabled because
dask-expr is not installed. [repeated 2x across cluster]
(_ray_fit pid=5693) You can install it with `pip install
dask[dataframe]` or `conda install dask`. [repeated 2x across cluster]
(_ray_fit pid=5693) This will raise in a future version. [repeated
2x across cluster]
(_ray_fit pid=5693) warnings.warn(msg, FutureWarning) [repeated
2x across cluster]
(_ray_fit pid=5693) [repeated 4x across cluster]
(_ray_fit pid=5817) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning: [repeated 2x across
cluster]
(_ray_fit pid=5817) Dask dataframe query planning is disabled because
dask-expr is not installed. [repeated 2x across cluster]
(_ray_fit pid=5817) You can install it with `pip install
dask[dataframe]` or `conda install dask`. [repeated 2x across cluster]

```

```

(_ray_fit pid=5817) This will raise in a future version. [repeated
2x across cluster]
(_ray_fit pid=5817) warnings.warn(msg, FutureWarning) [repeated
2x across cluster]
(_ray_fit pid=5817) [repeated 4x across cluster]
(_ray_fit pid=5896) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning: [repeated 2x across
cluster]
(_ray_fit pid=5896) Dask dataframe query planning is disabled because
dask-expr is not installed. [repeated 2x across cluster]
(_ray_fit pid=5896) You can install it with `pip install
dask[dataframe]` or `conda install dask`. [repeated 2x across cluster]
(_ray_fit pid=5896) This will raise in a future version. [repeated
2x across cluster]
(_ray_fit pid=5896) warnings.warn(msg, FutureWarning) [repeated
2x across cluster]
(_ray_fit pid=5896) [repeated 4x across cluster]
(_dystack pid=4697) -491.0254 = Validation score
(-root_mean_squared_error)
(_dystack pid=4697) 40.91s = Training runtime
(_dystack pid=4697) 1.16s = Validation runtime
(_ray_fit pid=5929) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning:
(_ray_fit pid=5929) Dask dataframe query planning is disabled because
dask-expr is not installed.
(_ray_fit pid=5929) You can install it with `pip install
dask[dataframe]` or `conda install dask`.
(_ray_fit pid=5929) This will raise in a future version.
(_ray_fit pid=5929) warnings.warn(msg, FutureWarning)
(_ray_fit pid=5929) [repeated 2x across cluster]
(_dystack pid=4697) Fitting model: WeightedEnsemble_L3 ... Training
model for up to 141.14s of the -34.97s of remaining time.
(_dystack pid=4697) Ensemble Weights: {'LightGBMXT_BAG_L1': 0.619,
'KNeighborsDist_BAG_L1': 0.19, 'LightGBMXT_BAG_L2': 0.19}
(_dystack pid=4697) -458.735 = Validation score
(-root_mean_squared_error)
(_dystack pid=4697) 0.02s = Training runtime
(_dystack pid=4697) 0.0s = Validation runtime
(_dystack pid=4697) AutoGluon training complete, total runtime =
176.43s ... Best model: WeightedEnsemble_L3 | Estimated inference throughput:
74.1 rows/s (2005 batch size)
(_dystack pid=4697) TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("agModels-predictprice/ds_sub_fit/sub_fit_ho")
(_dystack pid=4697) /usr/local/lib/python3.10/dist-
packages/dask/dataframe/__init__.py:42: FutureWarning:

```



```

(_dystack pid=4697) Dask dataframe query planning is disabled because
dask-expr is not installed.
(_dystack pid=4697) You can install it with `pip install
dask[dataframe]` or `conda install dask`.
(_dystack pid=4697) This will raise in a future version.
(_dystack pid=4697) warnings.warn(msg, FutureWarning)
(_dystack pid=4697) Deleting DyStack predictor artifacts
(clean_up_fits=True) ...
(_dystack pid=4697) [repeated 2x across cluster]
Leaderboard on holdout data (DyStack):

```

	model	score_holdout	score_val	eval_metric
0	WeightedEnsemble_L3	-327.077878	-458.734984	root_mean_squared_error
		27.591639	27.136705	158.261909
		0.000915	0.022242	3
				True
				6
1	WeightedEnsemble_L2	-330.105189	-460.711848	root_mean_squared_error
		26.656610	25.917874	117.333352
		0.000814	0.023995	2
				True
				4
2	LightGBMXT_BAG_L2	-345.434855	-491.025375	root_mean_squared_error
		27.588542	27.135790	158.239667
		1.162218	40.909981	2
				True
				5
3	LightGBMXT_BAG_L1	-352.453263	-484.984850	root_mean_squared_error
		26.635866	25.866909	117.284312
		25.866909	117.284312	1
				True
				3
4	KNeighborsDist_BAG_L1	-586.430480	-685.578353	root_mean_squared_error
		0.016799	0.050151	0.025044
		0.050151	0.025044	1
				True
				2
5	KNeighborsUnif_BAG_L1	-685.024589	-802.101549	root_mean_squared_error
		0.015565	0.056512	0.020329
		0.056512	0.020329	1
				True
				1

```

1 = Optimal num_stack_levels (Stacked Overfitting Occurred:
False)
216s = DyStack runtime | 384s = Remaining runtime
Starting main fit with num_stack_levels=1.
For future fit calls on this dataset, you can skip DyStack to save time:
`predictor.fit(..., dynamic_stacking=False, num_stack_levels=1)`
Beginning AutoGluon training ... Time limit = 384s
AutoGluon will save models to "agModels-predictprice"
Train Data Rows: 18039
Train Data Columns: 14
Label Column: price
Problem Type: regression
Preprocessing data ...
Using Feature Generators to preprocess the data ...
Fitting AutoMLPipelineFeatureGenerator...
Available Memory: 10178.14 MB
Train Data (Original) Memory Usage: 12.43 MB (0.1% of available memory)

```

```

    Inferring data type of each feature based on column values. Set
    feature_metadata_in to manually specify special dtypes of the features.
    Stage 1 Generators:
        Fitting AsTypeFeatureGenerator...
    Stage 2 Generators:
        Fitting FillNaFeatureGenerator...
    Stage 3 Generators:
        Fitting IdentityFeatureGenerator...
        Fitting CategoryFeatureGenerator...
        Fitting CategoryMemoryMinimizeFeatureGenerator...
    Stage 4 Generators:
        Fitting DropUniqueFeatureGenerator...
    Stage 5 Generators:
        Fitting DropDuplicatesFeatureGenerator...
    Types of features in original data (raw dtype, special dtypes):
        ('float', []) : 2 | ['latitude', 'longitude']
        ('object', []) : 12 | ['city', 'province', 'lease_term', 'type',
'beds', ...]
    Types of features in processed data (raw dtype, special dtypes):
        ('category', []) : 12 | ['city', 'province', 'lease_term',
'type', 'beds', ...]
        ('float', []) : 2 | ['latitude', 'longitude']
    0.5s = Fit runtime
    14 features in original data used to generate 14 features in processed
    data.
    Train Data (Processed) Memory Usage: 0.52 MB (0.0% of available memory)
    Data preprocessing and feature engineering runtime = 0.64s ...
    AutoGluon will gauge predictive performance using evaluation metric:
    'root_mean_squared_error'
    This metric's sign has been flipped to adhere to being higher_is_better.
    The metric score can be multiplied by -1 to get the metric value.
    To change this, specify the eval_metric parameter of Predictor()
    Large model count detected (112 configs) ... Only displaying the first 3 models
    of each family. To see all, set `verbosity=3`.
    User-specified model hyperparameters to be fit:
    {
        'NN_TORCH': [{}, {'activation': 'elu', 'dropout_prob':
0.10077639529843717, 'hidden_size': 108, 'learning_rate': 0.002735937344002146,
'num_layers': 4, 'use_batchnorm': True, 'weight_decay': 1.356433327634438e-12,
'ag_args': {'name_suffix': '_r79', 'priority': -2}}, {'activation': 'elu',
'dropout_prob': 0.11897478034205347, 'hidden_size': 213, 'learning_rate':
0.0010474382260641949, 'num_layers': 4, 'use_batchnorm': False, 'weight_decay':
5.594471067786272e-10, 'ag_args': {'name_suffix': '_r22', 'priority': -7}}],
        'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}], {},
'GBMLarge'],
        'CAT': [{}, {'depth': 6, 'grow_policy': 'SymmetricTree', 'l2_leaf_reg':
2.1542798306067823, 'learning_rate': 0.06864209415792857, 'max_ctr_complexity':
4, 'one_hot_max_size': 10, 'ag_args': {'name_suffix': '_r177', 'priority': -1}},

```



```
{'depth': 8, 'grow_policy': 'Depthwise', 'l2_leaf_reg': 2.7997999596449104,
'learning_rate': 0.031375015734637225, 'max_ctr_complexity': 2,
'one_hot_max_size': 3, 'ag_args': {'name_suffix': '_r9', 'priority': -5}},
  'XGB': [{}, {'colsample_bytree': 0.6917311125174739,
'enable_categorical': False, 'learning_rate': 0.018063876087523967, 'max_depth':
10, 'min_child_weight': 0.6028633586934382, 'ag_args': {'name_suffix': '_r33',
'priority': -8}}, {'colsample_bytree': 0.6628423832084077, 'enable_categorical':
False, 'learning_rate': 0.08775715546881824, 'max_depth': 5, 'min_child_weight':
0.6294123374222513, 'ag_args': {'name_suffix': '_r89', 'priority': -16}}],
  'FASTAI': [{}, {'bs': 256, 'emb_drop': 0.5411770367537934, 'epochs': 43,
'layers': [800, 400], 'lr': 0.01519848858318159, 'ps': 0.23782946566604385,
'ag_args': {'name_suffix': '_r191', 'priority': -4}}, {'bs': 2048, 'emb_drop':
0.05070411322605811, 'epochs': 29, 'layers': [200, 100], 'lr':
0.08974235041576624, 'ps': 0.10393466140748028, 'ag_args': {'name_suffix':
'_r102', 'priority': -11}}],
  'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
  'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
  'KNN': [{'weights': 'uniform', 'ag_args': {'name_suffix': 'Unif'}},
{'weights': 'distance', 'ag_args': {'name_suffix': 'Dist'}}],
}
```

AutoGluon will fit 2 stack levels (L1 to L2) ...

Fitting 108 L1 models ...

Fitting model: KNeighborsUnif\_BAG\_L1 ... Training model for up to 255.64s of the  
383.53s of remaining time.

-774.1865 = Validation score (-root\_mean\_squared\_error)

0.03s = Training runtime

0.09s = Validation runtime

Fitting model: KNeighborsDist\_BAG\_L1 ... Training model for up to 255.45s of the  
383.34s of remaining time.

-652.2715 = Validation score (-root\_mean\_squared\_error)

0.05s = Training runtime

0.09s = Validation runtime

Fitting model: LightGBMXT\_BAG\_L1 ... Training model for up to 255.27s of the  
383.15s of remaining time.

Fitting 8 child models (S1F1 - S1F8) | Fitting with

ParallelLocalFoldFittingStrategy (2 workers, per: cpus=1, gpus=0, memory=0.07%)

-442.0899 = Validation score (-root\_mean\_squared\_error)

282.43s = Training runtime

118.27s = Validation runtime

Fitting model: WeightedEnsemble\_L2 ... Training model for up to 360.0s of the

```

84.94s of remaining time.
    Ensemble Weights: {'LightGBMXT_BAG_L1': 0.792, 'KNeighborsDist_BAG_L1':
0.208}
    -423.1634          = Validation score    (-root_mean_squared_error)
    0.02s             = Training    runtime
    0.0s              = Validation runtime
Fitting 106 L2 models ...
Fitting model: LightGBMXT_BAG_L2 ... Training model for up to 84.89s of the
84.86s of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy (2 workers, per: cpus=1, gpus=0, memory=0.09%)
    -432.8841          = Validation score    (-root_mean_squared_error)
    63.47s            = Training    runtime
    12.98s            = Validation runtime
Fitting model: LightGBM_BAG_L2 ... Training model for up to 16.2s of the 16.16s
of remaining time.
    Fitting 8 child models (S1F1 - S1F8) | Fitting with
ParallelLocalFoldFittingStrategy (2 workers, per: cpus=1, gpus=0, memory=0.08%)
    -438.0843          = Validation score    (-root_mean_squared_error)
    29.97s            = Training    runtime
    1.08s             = Validation runtime
Fitting model: WeightedEnsemble_L3 ... Training model for up to 360.0s of the
-18.07s of remaining time.
    Ensemble Weights: {'LightGBMXT_BAG_L1': 0.4, 'LightGBMXT_BAG_L2': 0.3,
'LightGBM_BAG_L2': 0.2, 'KNeighborsDist_BAG_L1': 0.1}
    -414.8761          = Validation score    (-root_mean_squared_error)
    0.05s             = Training    runtime
    0.0s              = Validation runtime
AutoGluon training complete, total runtime = 402.39s ... Best model:
WeightedEnsemble_L3 | Estimated inference throughput: 17.0 rows/s (2255 batch
size)
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("agModels-predictprice")

```

```

[19]: test_data = TabularDataset(df.drop(train_data.index))
      y_test = test_data[target]
      # predictor = TabularPredictor.load("agModels-predictprice")

      y_pred = predictor_price.predict(test_data)
      print("Predictions: \n", y_pred)
      perf = predictor_price.evaluate_predictions(y_true=y_test, y_pred=y_pred,
      ↪auxiliary_metrics=True)

```

```

Predictions:
10      2818.560791
11      2026.548340
13      2583.564453
19      2126.927246

```

```

24      1997.721191
...
25748    1244.216064
25755    1066.665649
25762     896.792847
25763    1111.311279
25764     896.792847
Name: price, Length: 7732, dtype: float32

```

```
[20]: # Take a look at the performance of the model
perf
```

```
[20]: {'root_mean_squared_error': -431.76314186113444,
      'mean_squared_error': -186419.41066979812,
      'mean_absolute_error': -196.76120907033103,
      'r2': 0.8036366987519679,
      'pearsonr': 0.8978123558305657,
      'median_absolute_error': -116.48834228515625}
```

```
[21]: pd.DataFrame({'y_test': y_test, 'y_pred': y_pred})
```

```
[21]:
```

	y_test	y_pred
10	1930.0	2818.560791
11	1700.0	2026.548340
13	3150.0	2583.564453
19	2300.0	2126.927246
24	1910.0	1997.721191
...	...	...
25748	1305.0	1244.216064
25755	1085.0	1066.665649
25762	945.0	896.792847
25763	1025.0	1111.311279
25764	995.0	896.792847

```
[7732 rows x 2 columns]
```

```
[22]: # from google.colab import files
# files.download('/content/agModels-predictprice/models/WeightedEnsemble_L3')
```

```
[ ]: # # Reduce model size
# predictor_price.reduce_memory_size(remove_data=True, save_space=True)

# predictor_price.leaderboard(extra_info=True)
# predictor_price.delete_models(models_to_keep=predictor_price.
↳ leaderboard()['model'][:3])
# predictor_price.save("saved_model")
```

## 2 Explainable AI

```
[15]: predictor_price = TabularPredictor.load("agModels-predictprice")
```

```
[16]: predictor_price.info()
```

```
/usr/local/lib/python3.10/dist-packages/dask/dataframe/__init__.py:42:
```

```
FutureWarning:
```

```
Dask dataframe query planning is disabled because dask-expr is not installed.
```

```
You can install it with `pip install dask[dataframe]` or `conda install dask`.
```

```
This will raise in a future version.
```

```
warnings.warn(msg, FutureWarning)
```

```
[16]: {'path': 'agModels-predictprice',  
      'label': 'price',  
      'random_state': 0,  
      'version': '1.1.1',  
      'features': ['city',  
                  'province',  
                  'latitude',  
                  'longitude',  
                  'lease_term',  
                  'type',  
                  'beds',  
                  'baths',  
                  'sq_feet',  
                  'furnishing',  
                  'availability_date',  
                  'smoking',  
                  'cats',  
                  'dogs'],  
      'feature_metadata_in':  
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddd0fa8f0>,  
      'time_fit_preprocessing': 0.6379239559173584,  
      'time_fit_training': 401.7536509037018,  
      'time_fit_total': 402.39157485961914,  
      'time_limit': 384.20341539382935,  
      'time_train_start': 1731860728.1672926,  
      'num_rows_train': 18039,  
      'num_cols_train': 14,  
      'num_rows_val': None,  
      'num_classes': None,  
      'problem_type': 'regression',  
      'eval_metric': 'root_mean_squared_error',  
      'best_model': 'WeightedEnsemble_L3',
```

```

'best_model_score_val': -414.8760953952414,
'best_model_stack_level': 3,
'num_models_trained': 7,
'num_bag_folds': 8,
'max_stack_level': 3,
'max_core_stack_level': 2,
'model_info': {'KNeighborsUnif_BAG_L1': {'name': 'KNeighborsUnif_BAG_L1',
    'model_type': 'StackerEnsembleModel',
    'problem_type': 'regression',
    'eval_metric': 'root_mean_squared_error',
    'stopping_metric': 'root_mean_squared_error',
    'fit_time': 0.029515743255615234,
    'num_classes': None,
    'quantile_levels': None,
    'predict_time': 0.09212350845336914,
    'val_score': -774.186517755969,
    'hyperparameters': {'use_orig_features': True,
        'max_base_models': 25,
        'max_base_models_per_type': 5,
        'save_bag_folds': True,
        'use_child_oof': True},
    'hyperparameters_fit': {},
    'hyperparameters_nondefault': ['use_child_oof'],
    'ag_args_fit': {'max_memory_usage_ratio': 1.0,
        'max_time_limit_ratio': 1.0,
        'max_time_limit': None,
        'min_time_limit': 0,
        'valid_raw_types': None,
        'valid_special_types': None,
        'ignored_type_group_special': None,
        'ignored_type_group_raw': None,
        'get_features_kwargs': None,
        'get_features_kwargs_extra': None,
        'predict_1_batch_size': None,
        'temperature_scalar': None,
        'drop_unique': False},
    'num_features': 2,
    'features': ['longitude', 'latitude'],
    'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddd12e290>,
    'memory_size': 2638,
    'compile_time': None,
    'is_initialized': True,
    'is_fit': True,
    'is_valid': True,
    'can_infer': True,
    'bagged_info': {'child_model_type': 'KNNModel',

```

```

'num_child_models': 1,
'child_model_names': ['S1F1'],
'n_repeats': 1,
'k_per_n_repeat': [1],
'_random_state': 1,
'low_memory': True,
'bagged_mode': False,
'max_memory_size': 792488,
'min_memory_size': 792488,
'child_hyperparameters': {'weights': 'uniform'},
'child_hyperparameters_fit': {},
'child_ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['int', 'float'],
'valid_special_types': None,
'ignored_type_group_special': ['bool'],
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None}},
'stacker_info': {'num_base_models': 0, 'base_model_names': []},
'children_info': {'S1F1': {'name': 'S1F1',
'model_type': 'KNNModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 0.029515743255615234,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.09212350845336914,
'val_score': -774.186517755969,
'hyperparameters': {'weights': 'uniform'},
'hyperparameters_fit': {},
'hyperparameters_nondefault': ['weights'],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['int', 'float'],
'valid_special_types': None,
'ignored_type_group_special': ['bool'],
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,

```

```

        'predict_1_batch_size': None,
        'temperature_scalar': None},
        'num_features': 2,
        'features': ['latitude', 'longitude'],
        'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aede741abc0>,
        'memory_size': 789850,
        'compile_time': None,
        'is_initialized': True,
        'is_fit': True,
        'is_valid': True,
        'can_infer': True}}},
'KNeighborsDist_BAG_L1': {'name': 'KNeighborsDist_BAG_L1',
        'model_type': 'StackerEnsembleModel',
        'problem_type': 'regression',
        'eval_metric': 'root_mean_squared_error',
        'stopping_metric': 'root_mean_squared_error',
        'fit_time': 0.04554128646850586,
        'num_classes': None,
        'quantile_levels': None,
        'predict_time': 0.0915365219116211,
        'val_score': -652.2714801027624,
        'hyperparameters': {'use_orig_features': True,
        'max_base_models': 25,
        'max_base_models_per_type': 5,
        'save_bag_folds': True,
        'use_child_oof': True},
        'hyperparameters_fit': {},
        'hyperparameters_nondefault': ['use_child_oof'],
        'ag_args_fit': {'max_memory_usage_ratio': 1.0,
        'max_time_limit_ratio': 1.0,
        'max_time_limit': None,
        'min_time_limit': 0,
        'valid_raw_types': None,
        'valid_special_types': None,
        'ignored_type_group_special': None,
        'ignored_type_group_raw': None,
        'get_features_kwargs': None,
        'get_features_kwargs_extra': None,
        'predict_1_batch_size': None,
        'temperature_scalar': None,
        'drop_unique': False},
        'num_features': 2,
        'features': ['longitude', 'latitude'],
        'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddd0fba90>,
        'memory_size': 2638,

```

```

'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True,
'bagged_info': {'child_model_type': 'KNNModel',
'num_child_models': 1,
'child_model_names': ['S1F1'],
'_n_repeats': 1,
'_k_per_n_repeat': [1],
'_random_state': 1,
'low_memory': True,
'bagged_mode': False,
'max_memory_size': 792489,
'min_memory_size': 792489,
'child_hyperparameters': {'weights': 'distance'},
'child_hyperparameters_fit': {},
'child_ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['int', 'float'],
'valid_special_types': None,
'ignored_type_group_special': ['bool'],
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None}},
'stacker_info': {'num_base_models': 0, 'base_model_names': []},
'children_info': {'S1F1': {'name': 'S1F1',
'model_type': 'KNNModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 0.04554128646850586,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.0915365219116211,
'val_score': -652.2714801027624,
'hyperparameters': {'weights': 'distance'},
'hyperparameters_fit': {},
'hyperparameters_nondefault': ['weights'],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,

```



```

    'valid_raw_types': ['int', 'float'],
    'valid_special_types': None,
    'ignored_type_group_special': ['bool'],
    'ignored_type_group_raw': None,
    'get_features_kwargs': None,
    'get_features_kwargs_extra': None,
    'predict_1_batch_size': None,
    'temperature_scalar': None},
    'num_features': 2,
    'features': ['latitude', 'longitude'],
    'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aee2435ac50>,
    'memory_size': 789851,
    'compile_time': None,
    'is_initialized': True,
    'is_fit': True,
    'is_valid': True,
    'can_infer': True}}},
'LightGBMXT_BAG_L1': {'name': 'LightGBMXT_BAG_L1',
    'model_type': 'StackerEnsembleModel',
    'problem_type': 'regression',
    'eval_metric': 'root_mean_squared_error',
    'stopping_metric': 'root_mean_squared_error',
    'fit_time': 282.42990374565125,
    'num_classes': None,
    'quantile_levels': None,
    'predict_time': 118.26582980155945,
    'val_score': -442.0899321234507,
    'hyperparameters': {'use_orig_features': True,
        'max_base_models': 25,
        'max_base_models_per_type': 5,
        'save_bag_folds': True},
    'hyperparameters_fit': {},
    'hyperparameters_nondefault': [],
    'ag_args_fit': {'max_memory_usage_ratio': 1.0,
        'max_time_limit_ratio': 1.0,
        'max_time_limit': None,
        'min_time_limit': 0,
        'valid_raw_types': None,
        'valid_special_types': None,
        'ignored_type_group_special': None,
        'ignored_type_group_raw': None,
        'get_features_kwargs': None,
        'get_features_kwargs_extra': None,
        'predict_1_batch_size': None,
        'temperature_scalar': None,
        'drop_unique': False},

```

```

'num_features': 14,
'features': ['province',
             'type',
             'city',
             'longitude',
             'cats',
             'beds',
             'smoking',
             'furnishing',
             'baths',
             'sq_feet',
             'dogs',
             'lease_term',
             'latitude',
             'availability_date'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddd0fb1c0>,
'memory_size': 2870,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True,
'bagged_info': {'child_model_type': 'LGBModel',
                 'num_child_models': 8,
                 'child_model_names': ['S1F2',
                                       'S1F1',
                                       'S1F4',
                                       'S1F3',
                                       'S1F5',
                                       'S1F6',
                                       'S1F7',
                                       'S1F8'],
                 '_n_repeats': 1,
                 '_k_per_n_repeat': [8],
                 '_random_state': 1,
                 'low_memory': True,
                 'bagged_mode': True,
                 'max_memory_size': 181928391,
                 'min_memory_size': 26095574,
                 'child_hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
                 'child_hyperparameters_fit': {'num_boost_round': 8712},
                 'child_ag_args_fit': {'max_memory_usage_ratio': 1.0,
                                       'max_time_limit_ratio': 1.0,
                                       'max_time_limit': None,
                                       'min_time_limit': 0,
                                       'valid_raw_types': ['bool', 'int', 'float', 'category']},

```

```

'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None}},
'stacker_info': {'num_base_models': 0, 'base_model_names': []},
'children_info': {'S1F2': {'name': 'S1F2',
    'model_type': 'LGBModel',
    'problem_type': 'regression',
    'eval_metric': 'root_mean_squared_error',
    'stopping_metric': 'root_mean_squared_error',
    'fit_time': 58.19691824913025,
    'num_classes': None,
    'quantile_levels': None,
    'predict_time': 20.48674750328064,
    'val_score': -424.3561316043692,
    'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
    'hyperparameters_fit': {'num_boost_round': 9645},
    'hyperparameters_nondefault': ['extra_trees'],
    'ag_args_fit': {'max_memory_usage_ratio': 1.0,
        'max_time_limit_ratio': 1.0,
        'max_time_limit': None,
        'min_time_limit': 0,
        'valid_raw_types': ['bool', 'int', 'float', 'category'],
        'valid_special_types': None,
        'ignored_type_group_special': None,
        'ignored_type_group_raw': None,
        'get_features_kwargs': None,
        'get_features_kwargs_extra': None,
        'predict_1_batch_size': None,
        'temperature_scalar': None},
    'num_features': 14,
    'features': ['latitude',
        'longitude',
        'city',
        'province',
        'lease_term',
        'type',
        'beds',
        'baths',
        'sq_feet',
        'furnishing',
        'availability_date',
        'smoking',
        'cats',

```

```

    'dogs'],
    'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aede741a560>,
    'memory_size': 25143815,
    'compile_time': None,
    'is_initialized': True,
    'is_fit': True,
    'is_valid': True,
    'can_infer': True},
    'S1F1': {'name': 'S1F1',
    'model_type': 'LGBModel',
    'problem_type': 'regression',
    'eval_metric': 'root_mean_squared_error',
    'stopping_metric': 'root_mean_squared_error',
    'fit_time': 59.269981145858765,
    'num_classes': None,
    'quantile_levels': None,
    'predict_time': 21.030065536499023,
    'val_score': -382.94466320599656,
    'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
    'hyperparameters_fit': {'num_boost_round': 9587},
    'hyperparameters_nondefault': ['extra_trees'],
    'ag_args_fit': {'max_memory_usage_ratio': 1.0,
    'max_time_limit_ratio': 1.0,
    'max_time_limit': None,
    'min_time_limit': 0,
    'valid_raw_types': ['bool', 'int', 'float', 'category'],
    'valid_special_types': None,
    'ignored_type_group_special': None,
    'ignored_type_group_raw': None,
    'get_features_kwargs': None,
    'get_features_kwargs_extra': None,
    'predict_1_batch_size': None,
    'temperature_scalar': None},
    'num_features': 14,
    'features': ['latitude',
    'longitude',
    'city',
    'province',
    'lease_term',
    'type',
    'beds',
    'baths',
    'sq_feet',
    'furnishing',
    'availability_date',
    'smoking',

```

```

        'cats',
        'dogs'],
        'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aee16875270>,
        'memory_size': 25018284,
        'compile_time': None,
        'is_initialized': True,
        'is_fit': True,
        'is_valid': True,
        'can_infer': True},
        'S1F4': {'name': 'S1F4',
        'model_type': 'LGBModel',
        'problem_type': 'regression',
        'eval_metric': 'root_mean_squared_error',
        'stopping_metric': 'root_mean_squared_error',
        'fit_time': 55.36878061294556,
        'num_classes': None,
        'quantile_levels': None,
        'predict_time': 16.661664962768555,
        'val_score': -368.63815774600977,
        'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
        'hyperparameters_fit': {'num_boost_round': 9935},
        'hyperparameters_nondefault': ['extra_trees'],
        'ag_args_fit': {'max_memory_usage_ratio': 1.0,
        'max_time_limit_ratio': 1.0,
        'max_time_limit': None,
        'min_time_limit': 0,
        'valid_raw_types': ['bool', 'int', 'float', 'category'],
        'valid_special_types': None,
        'ignored_type_group_special': None,
        'ignored_type_group_raw': None,
        'get_features_kwargs': None,
        'get_features_kwargs_extra': None,
        'predict_1_batch_size': None,
        'temperature_scalar': None},
        'num_features': 14,
        'features': ['latitude',
        'longitude',
        'city',
        'province',
        'lease_term',
        'type',
        'beds',
        'baths',
        'sq_feet',
        'furnishing',
        'availability_date',

```

```

        'smoking',
        'cats',
        'dogs'],
        'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddcbc4100>,
        'memory_size': 25916132,
        'compile_time': None,
        'is_initialized': True,
        'is_fit': True,
        'is_valid': True,
        'can_infer': True},
        'S1F3': {'name': 'S1F3',
        'model_type': 'LGBModel',
        'problem_type': 'regression',
        'eval_metric': 'root_mean_squared_error',
        'stopping_metric': 'root_mean_squared_error',
        'fit_time': 55.31507325172424,
        'num_classes': None,
        'quantile_levels': None,
        'predict_time': 16.58866047859192,
        'val_score': -649.5601056189936,
        'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
        'hyperparameters_fit': {'num_boost_round': 9887},
        'hyperparameters_nondefault': ['extra_trees'],
        'ag_args_fit': {'max_memory_usage_ratio': 1.0,
        'max_time_limit_ratio': 1.0,
        'max_time_limit': None,
        'min_time_limit': 0,
        'valid_raw_types': ['bool', 'int', 'float', 'category'],
        'valid_special_types': None,
        'ignored_type_group_special': None,
        'ignored_type_group_raw': None,
        'get_features_kwargs': None,
        'get_features_kwargs_extra': None,
        'predict_1_batch_size': None,
        'temperature_scalar': None},
        'num_features': 14,
        'features': ['latitude',
        'longitude',
        'city',
        'province',
        'lease_term',
        'type',
        'beds',
        'baths',
        'sq_feet',
        'furnishing',

```

```

    'availability_date',
    'smoking',
    'cats',
    'dogs'],
    'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddd0fb070>,
    'memory_size': 25869183,
    'compile_time': None,
    'is_initialized': True,
    'is_fit': True,
    'is_valid': True,
    'can_infer': True},
    'S1F5': {'name': 'S1F5',
    'model_type': 'LGBModel',
    'problem_type': 'regression',
    'eval_metric': 'root_mean_squared_error',
    'stopping_metric': 'root_mean_squared_error',
    'fit_time': 52.14553451538086,
    'num_classes': None,
    'quantile_levels': None,
    'predict_time': 16.65358066558838,
    'val_score': -417.697761743281,
    'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
    'hyperparameters_fit': {'num_boost_round': 9992},
    'hyperparameters_nondefault': ['extra_trees'],
    'ag_args_fit': {'max_memory_usage_ratio': 1.0,
    'max_time_limit_ratio': 1.0,
    'max_time_limit': None,
    'min_time_limit': 0,
    'valid_raw_types': ['bool', 'int', 'float', 'category'],
    'valid_special_types': None,
    'ignored_type_group_special': None,
    'ignored_type_group_raw': None,
    'get_features_kwargs': None,
    'get_features_kwargs_extra': None,
    'predict_1_batch_size': None,
    'temperature_scalar': None},
    'num_features': 14,
    'features': ['latitude',
    'longitude',
    'city',
    'province',
    'lease_term',
    'type',
    'beds',
    'baths',
    'sq_feet',

```

```

        'furnishing',
        'availability_date',
        'smoking',
        'cats',
        'dogs'],
        'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddcd09540>,
        'memory_size': 26064503,
        'compile_time': None,
        'is_initialized': True,
        'is_fit': True,
        'is_valid': True,
        'can_infer': True},
        'S1F6': {'name': 'S1F6',
        'model_type': 'LGBModel',
        'problem_type': 'regression',
        'eval_metric': 'root_mean_squared_error',
        'stopping_metric': 'root_mean_squared_error',
        'fit_time': 52.45994424819946,
        'num_classes': None,
        'quantile_levels': None,
        'predict_time': 16.317374229431152,
        'val_score': -424.5945086276661,
        'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
        'hyperparameters_fit': {'num_boost_round': 9942},
        'hyperparameters_nondefault': ['extra_trees'],
        'ag_args_fit': {'max_memory_usage_ratio': 1.0,
        'max_time_limit_ratio': 1.0,
        'max_time_limit': None,
        'min_time_limit': 0,
        'valid_raw_types': ['bool', 'int', 'float', 'category'],
        'valid_special_types': None,
        'ignored_type_group_special': None,
        'ignored_type_group_raw': None,
        'get_features_kwargs': None,
        'get_features_kwargs_extra': None,
        'predict_1_batch_size': None,
        'temperature_scalar': None},
        'num_features': 14,
        'features': ['latitude',
        'longitude',
        'city',
        'province',
        'lease_term',
        'type',
        'beds',
        'baths',

```



```

        'sq_feet',
        'furnishing',
        'availability_date',
        'smoking',
        'cats',
        'dogs'],
        'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddd0fad40>,
        'memory_size': 25938441,
        'compile_time': None,
        'is_initialized': True,
        'is_fit': True,
        'is_valid': True,
        'can_infer': True},
        'S1F7': {'name': 'S1F7',
        'model_type': 'LGBModel',
        'problem_type': 'regression',
        'eval_metric': 'root_mean_squared_error',
        'stopping_metric': 'root_mean_squared_error',
        'fit_time': 5.27716851234436,
        'num_classes': None,
        'quantile_levels': None,
        'predict_time': 0.4103834629058838,
        'val_score': -411.86323027391967,
        'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
        'hyperparameters_fit': {'num_boost_round': 707},
        'hyperparameters_nondefault': ['extra_trees'],
        'ag_args_fit': {'max_memory_usage_ratio': 1.0,
        'max_time_limit_ratio': 1.0,
        'max_time_limit': None,
        'min_time_limit': 0,
        'valid_raw_types': ['bool', 'int', 'float', 'category'],
        'valid_special_types': None,
        'ignored_type_group_special': None,
        'ignored_type_group_raw': None,
        'get_features_kwargs': None,
        'get_features_kwargs_extra': None,
        'predict_1_batch_size': None,
        'temperature_scalar': None},
        'num_features': 14,
        'features': ['latitude',
        'longitude',
        'city',
        'province',
        'lease_term',
        'type',
        'beds',

```

```

    'baths',
    'sq_feet',
    'furnishing',
    'availability_date',
    'smoking',
    'cats',
    'dogs'],
    'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddd0fa590>,
    'memory_size': 1882459,
    'compile_time': None,
    'is_initialized': True,
    'is_fit': True,
    'is_valid': True,
    'can_infer': True},
    'S1F8': {'name': 'S1F8',
    'model_type': 'LGBModel',
    'problem_type': 'regression',
    'eval_metric': 'root_mean_squared_error',
    'stopping_metric': 'root_mean_squared_error',
    'fit_time': 31.96611976623535,
    'num_classes': None,
    'quantile_levels': None,
    'predict_time': 10.117352962493896,
    'val_score': -393.1914940836134,
    'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
    'hyperparameters_fit': {'num_boost_round': 10000},
    'hyperparameters_nondefault': ['extra_trees'],
    'ag_args_fit': {'max_memory_usage_ratio': 1.0,
    'max_time_limit_ratio': 1.0,
    'max_time_limit': None,
    'min_time_limit': 0,
    'valid_raw_types': ['bool', 'int', 'float', 'category'],
    'valid_special_types': None,
    'ignored_type_group_special': None,
    'ignored_type_group_raw': None,
    'get_features_kwargs': None,
    'get_features_kwargs_extra': None,
    'predict_1_batch_size': None,
    'temperature_scalar': None},
    'num_features': 14,
    'features': ['latitude',
    'longitude',
    'city',
    'province',
    'lease_term',
    'type',

```

```

        'beds',
        'baths',
        'sq_feet',
        'furnishing',
        'availability_date',
        'smoking',
        'cats',
        'dogs'],
        'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddcd096f0>,
        'memory_size': 26092704,
        'compile_time': None,
        'is_initialized': True,
        'is_fit': True,
        'is_valid': True,
        'can_infer': True}}},
'WeightedEnsemble_L2': {'name': 'WeightedEnsemble_L2',
        'model_type': 'WeightedEnsembleModel',
        'problem_type': 'regression',
        'eval_metric': 'root_mean_squared_error',
        'stopping_metric': 'root_mean_squared_error',
        'fit_time': 0.023722410202026367,
        'num_classes': None,
        'quantile_levels': None,
        'predict_time': 0.0008921623229980469,
        'val_score': -423.163408198703,
        'hyperparameters': {'use_orig_features': False,
        'max_base_models': 25,
        'max_base_models_per_type': 5,
        'save_bag_folds': True},
        'hyperparameters_fit': {},
        'hyperparameters_nondefault': ['save_bag_folds'],
        'ag_args_fit': {'max_memory_usage_ratio': 1.0,
        'max_time_limit_ratio': 1.0,
        'max_time_limit': None,
        'min_time_limit': 0,
        'valid_raw_types': None,
        'valid_special_types': None,
        'ignored_type_group_special': None,
        'ignored_type_group_raw': None,
        'get_features_kwargs': None,
        'get_features_kwargs_extra': None,
        'predict_1_batch_size': None,
        'temperature_scalar': None,
        'drop_unique': False},
        'num_features': 2,
        'features': ['LightGBMXT_BAG_L1', 'KNeighborsDist_BAG_L1'],

```

```

'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aedd751e10>,
'memory_size': 2963,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True,
'bagged_info': {'child_model_type': 'GreedyWeightedEnsembleModel',
'num_child_models': 1,
'child_model_names': ['S1F1'],
'_n_repeats': 1,
'_k_per_n_repeat': [1],
'_random_state': 2,
'low_memory': False,
'bagged_mode': False,
'max_memory_size': 2963,
'min_memory_size': 2963,
'child_hyperparameters': {'ensemble_size': 25, 'subsample_size': 1000000},
'child_hyperparameters_fit': {'ensemble_size': 24},
'child_ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': None,
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None,
'drop_unique': False}},
'stacker_info': {'num_base_models': 2,
'base_model_names': ['KNeighborsDist_BAG_L1', 'LightGBMXT_BAG_L1']},
'children_info': {'S1F1': {'name': 'S1F1',
'model_type': 'GreedyWeightedEnsembleModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 0.023722410202026367,
'num_classes': None,
'quantile_levels': None,
'predict_time': None,
'val_score': None,
'hyperparameters': {'ensemble_size': 25, 'subsample_size': 1000000},
'hyperparameters_fit': {'ensemble_size': 24},

```

```

'hyperparameters_nondefault': [],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': None,
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None,
'drop_unique': False},
'num_features': 2,
'features': ['KNeighborsDist_BAG_L1', 'LightGBMXT_BAG_L1'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aedd751cf0>,
'memory_size': 5956,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True,
'model_weights': {'KNeighborsDist_BAG_L1': 0.20833333333333334,
'LightGBMXT_BAG_L1': 0.7916666666666666}}},
'LightGBMXT_BAG_L2': {'name': 'LightGBMXT_BAG_L2',
'model_type': 'StackerEnsembleModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 63.47079634666443,
'num_classes': None,
'quantile_levels': None,
'predict_time': 12.97685432434082,
'val_score': -432.88414922328326,
'hyperparameters': {'use_orig_features': True,
'max_base_models': 25,
'max_base_models_per_type': 5,
'save_bag_folds': True},
'hyperparameters_fit': {},
'hyperparameters_nondefault': [],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': None,

```

```

'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None,
'drop_unique': False},
'num_features': 17,
'features': ['province',
'LightGBMXT_BAG_L1',
'type',
'city',
'KNeighborsUnif_BAG_L1',
'longitude',
'cats',
'beds',
'smoking',
'KNeighborsDist_BAG_L1',
'furnishing',
'baths',
'sq_feet',
'dogs',
'lease_term',
'latitude',
'availability_date'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddcd1c220>,
'memory_size': 3208,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True,
'bagged_info': {'child_model_type': 'LGBModel',
'num_child_models': 8,
'child_model_names': ['S1F1',
'S1F3',
'S1F2',
'S1F4',
'S1F5',
'S1F6',
'S1F7',
'S1F8']},
'_n_repeats': 1,
'_k_per_n_repeat': [8],
'_random_state': 2,

```

```

'low_memory': True,
'bagged_mode': True,
'max_memory_size': 31486982,
'min_memory_size': 9479445,
'child_hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
'child_hyperparameters_fit': {'num_boost_round': 1479},
'child_ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None}},
'stacker_info': {'num_base_models': 3,
'base_model_names': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1']},
'children_info': {'S1F1': {'name': 'S1F1',
'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 12.375351667404175,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.6218945980072021,
'val_score': -396.1039605761208,
'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
'hyperparameters_fit': {'num_boost_round': 1143},
'hyperparameters_nondefault': ['extra_trees'],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},

```

```

'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
             'KNeighborsDist_BAG_L1',
             'LightGBMXT_BAG_L1',
             'latitude',
             'longitude',
             'city',
             'province',
             'lease_term',
             'type',
             'beds',
             'baths',
             'sq_feet',
             'furnishing',
             'availability_date',
             'smoking',
             'cats',
             'dogs'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aee24385240>,
'memory_size': 3039448,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True},
'S1F3': {'name': 'S1F3',
'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 2.643284797668457,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.03850364685058594,
'val_score': -324.3544768088787,
'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
'hyperparameters_fit': {'num_boost_round': 93},
'hyperparameters_nondefault': ['extra_trees'],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,

```



```

'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},
'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aee24387b20>,
'memory_size': 272574,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True},
'S1F2': {'name': 'S1F2',
'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 22.0169997215271,
'num_classes': None,
'quantile_levels': None,
'predict_time': 3.1859657764434814,
'val_score': -659.9561650108415,
'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
'hyperparameters_fit': {'num_boost_round': 3099},
'hyperparameters_nondefault': ['extra_trees'],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,

```

```

'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},
'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddd0fb5e0>,
'memory_size': 8211234,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True},
'S1F4': {'name': 'S1F4',
'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 3.354447364807129,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.059639930725097656,
'val_score': -380.04525872457435,
'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
'hyperparameters_fit': {'num_boost_round': 137},
'hyperparameters_nondefault': ['extra_trees'],

```

```

'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},
'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aee243876d0>,
'memory_size': 388929,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True},
'S1F5': {'name': 'S1F5',
'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 19.34431219100952,
'num_classes': None,
'quantile_levels': None,
'predict_time': 3.263979434967041,

```

```

'val_score': -417.26337347013185,
'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
'hyperparameters_fit': {'num_boost_round': 3577},
'hyperparameters_nondefault': ['extra_trees'],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},
'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aee24387220>,
'memory_size': 9476237,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True},
'S1F6': {'name': 'S1F6',
'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',

```

```

'fit_time': 19.204747676849365,
'num_classes': None,
'quantile_levels': None,
'predict_time': 5.70510458946228,
'val_score': -426.7062673536995,
'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
'hyperparameters_fit': {'num_boost_round': 3514},
'hyperparameters_nondefault': ['extra_trees'],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},
'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aee24384490>,
'memory_size': 9328076,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True},
'S1F7': {'name': 'S1F7',

```

```

'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 3.099177122116089,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.07348895072937012,
'val_score': -413.83553159023853,
'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
'hyperparameters_fit': {'num_boost_round': 182},
'hyperparameters_nondefault': ['extra_trees'],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},
'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aee24359390>,
'memory_size': 508720,
'compile_time': None,
'is_initialized': True,

```

```

'is_fit': True,
'is_valid': True,
'can_infer': True},
'S1F8': {'name': 'S1F8',
'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 2.3083550930023193,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.02827739715576172,
'val_score': -359.93933605571164,
'hyperparameters': {'learning_rate': 0.05, 'extra_trees': True},
'hyperparameters_fit': {'num_boost_round': 88},
'hyperparameters_nondefault': ['extra_trees'],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},
'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':

```

```

<autoglulon.common.features.feature_metadata.FeatureMetadata at 0x7aeddcd09630>,
  'memory_size': 258556,
  'compile_time': None,
  'is_initialized': True,
  'is_fit': True,
  'is_valid': True,
  'can_infer': True}}},
'LightGBM_BAG_L2': {'name': 'LightGBM_BAG_L2',
  'model_type': 'StackerEnsembleModel',
  'problem_type': 'regression',
  'eval_metric': 'root_mean_squared_error',
  'stopping_metric': 'root_mean_squared_error',
  'fit_time': 29.96975064277649,
  'num_classes': None,
  'quantile_levels': None,
  'predict_time': 1.0781099796295166,
  'val_score': -438.08432384123813,
  'hyperparameters': {'use_orig_features': True,
    'max_base_models': 25,
    'max_base_models_per_type': 5,
    'save_bag_folds': True},
  'hyperparameters_fit': {},
  'hyperparameters_nondefault': [],
  'ag_args_fit': {'max_memory_usage_ratio': 1.0,
    'max_time_limit_ratio': 1.0,
    'max_time_limit': None,
    'min_time_limit': 0,
    'valid_raw_types': None,
    'valid_special_types': None,
    'ignored_type_group_special': None,
    'ignored_type_group_raw': None,
    'get_features_kwargs': None,
    'get_features_kwargs_extra': None,
    'predict_1_batch_size': None,
    'temperature_scalar': None,
    'drop_unique': False},
  'num_features': 17,
  'features': ['province',
    'LightGBMXT_BAG_L1',
    'type',
    'city',
    'KNeighborsUnif_BAG_L1',
    'longitude',
    'cats',
    'beds',
    'smoking',
    'KNeighborsDist_BAG_L1',

```



```

'furnishing',
'baths',
'sq_feet',
'dogs',
'lease_term',
'latitude',
'availability_date'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddd0fb160>,
'memory_size': 3204,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True,
'bagged_info': {'child_model_type': 'LGBModel',
'num_child_models': 8,
'child_model_names': ['S1F1',
'S1F2',
'S1F3',
'S1F4',
'S1F5',
'S1F6',
'S1F7',
'S1F8'],
'_n_repeats': 1,
'_k_per_n_repeat': [8],
'_random_state': 2,
'low_memory': True,
'bagged_mode': True,
'max_memory_size': 5045688,
'min_memory_size': 1356557,
'child_hyperparameters': {'learning_rate': 0.05},
'child_hyperparameters_fit': {'num_boost_round': 202},
'child_ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None}},
'stacker_info': {'num_base_models': 3,

```

```

'base_model_names': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1']},
'children_info': {'S1F1': {'name': 'S1F1',
'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 5.02400279045105,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.05747342109680176,
'val_score': -364.9733263439527,
'hyperparameters': {'learning_rate': 0.05},
'hyperparameters_fit': {'num_boost_round': 128},
'hyperparameters_nondefault': [],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},
'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':

```

```

<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aee2435a290>,
  'memory_size': 414331,
  'compile_time': None,
  'is_initialized': True,
  'is_fit': True,
  'is_valid': True,
  'can_infer': True},
'S1F2': {'name': 'S1F2',
  'model_type': 'LGBModel',
  'problem_type': 'regression',
  'eval_metric': 'root_mean_squared_error',
  'stopping_metric': 'root_mean_squared_error',
  'fit_time': 5.455013751983643,
  'num_classes': None,
  'quantile_levels': None,
  'predict_time': 0.2831716537475586,
  'val_score': -703.3882333580328,
  'hyperparameters': {'learning_rate': 0.05},
  'hyperparameters_fit': {'num_boost_round': 355},
  'hyperparameters_nondefault': [],
  'ag_args_fit': {'max_memory_usage_ratio': 1.0,
    'max_time_limit_ratio': 1.0,
    'max_time_limit': None,
    'min_time_limit': 0,
    'valid_raw_types': ['bool', 'int', 'float', 'category'],
    'valid_special_types': None,
    'ignored_type_group_special': None,
    'ignored_type_group_raw': None,
    'get_features_kwargs': None,
    'get_features_kwargs_extra': None,
    'predict_1_batch_size': None,
    'temperature_scalar': None},
  'num_features': 17,
  'features': ['KNeighborsUnif_BAG_L1',
    'KNeighborsDist_BAG_L1',
    'LightGBMXT_BAG_L1',
    'latitude',
    'longitude',
    'city',
    'province',
    'lease_term',
    'type',
    'beds',
    'baths',
    'sq_feet',
    'furnishing',
    'availability_date',

```

```

    'smoking',
    'cats',
    'dogs'],
    'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddcd096c0>,
    'memory_size': 1103362,
    'compile_time': None,
    'is_initialized': True,
    'is_fit': True,
    'is_valid': True,
    'can_infer': True},
    'S1F3': {'name': 'S1F3',
    'model_type': 'LGBModel',
    'problem_type': 'regression',
    'eval_metric': 'root_mean_squared_error',
    'stopping_metric': 'root_mean_squared_error',
    'fit_time': 2.7538435459136963,
    'num_classes': None,
    'quantile_levels': None,
    'predict_time': 0.022531986236572266,
    'val_score': -347.70574380996885,
    'hyperparameters': {'learning_rate': 0.05},
    'hyperparameters_fit': {'num_boost_round': 49},
    'hyperparameters_nondefault': [],
    'ag_args_fit': {'max_memory_usage_ratio': 1.0,
    'max_time_limit_ratio': 1.0,
    'max_time_limit': None,
    'min_time_limit': 0,
    'valid_raw_types': ['bool', 'int', 'float', 'category'],
    'valid_special_types': None,
    'ignored_type_group_special': None,
    'ignored_type_group_raw': None,
    'get_features_kwargs': None,
    'get_features_kwargs_extra': None,
    'predict_1_batch_size': None,
    'temperature_scalar': None},
    'num_features': 17,
    'features': ['KNeighborsUnif_BAG_L1',
    'KNeighborsDist_BAG_L1',
    'LightGBMXT_BAG_L1',
    'latitude',
    'longitude',
    'city',
    'province',
    'lease_term',
    'type',
    'beds',

```

```

    'baths',
    'sq_feet',
    'furnishing',
    'availability_date',
    'smoking',
    'cats',
    'dogs'],
    'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddccf55900>,
    'memory_size': 161804,
    'compile_time': None,
    'is_initialized': True,
    'is_fit': True,
    'is_valid': True,
    'can_infer': True},
    'S1F4': {'name': 'S1F4',
    'model_type': 'LGBModel',
    'problem_type': 'regression',
    'eval_metric': 'root_mean_squared_error',
    'stopping_metric': 'root_mean_squared_error',
    'fit_time': 3.5004281997680664,
    'num_classes': None,
    'quantile_levels': None,
    'predict_time': 0.09077143669128418,
    'val_score': -365.3430320179888,
    'hyperparameters': {'learning_rate': 0.05},
    'hyperparameters_fit': {'num_boost_round': 115},
    'hyperparameters_nondefault': [],
    'ag_args_fit': {'max_memory_usage_ratio': 1.0,
    'max_time_limit_ratio': 1.0,
    'max_time_limit': None,
    'min_time_limit': 0,
    'valid_raw_types': ['bool', 'int', 'float', 'category'],
    'valid_special_types': None,
    'ignored_type_group_special': None,
    'ignored_type_group_raw': None,
    'get_features_kwargs': None,
    'get_features_kwargs_extra': None,
    'predict_1_batch_size': None,
    'temperature_scalar': None},
    'num_features': 17,
    'features': ['KNeighborsUnif_BAG_L1',
    'KNeighborsDist_BAG_L1',
    'LightGBMXT_BAG_L1',
    'latitude',
    'longitude',
    'city',

```

```

    'province',
    'lease_term',
    'type',
    'beds',
    'baths',
    'sq_feet',
    'furnishing',
    'availability_date',
    'smoking',
    'cats',
    'dogs'],
    'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddcf556c0>,
    'memory_size': 374502,
    'compile_time': None,
    'is_initialized': True,
    'is_fit': True,
    'is_valid': True,
    'can_infer': True},
    'S1F5': {'name': 'S1F5',
    'model_type': 'LGBModel',
    'problem_type': 'regression',
    'eval_metric': 'root_mean_squared_error',
    'stopping_metric': 'root_mean_squared_error',
    'fit_time': 5.55760645866394,
    'num_classes': None,
    'quantile_levels': None,
    'predict_time': 0.21150636672973633,
    'val_score': -433.28387639263207,
    'hyperparameters': {'learning_rate': 0.05},
    'hyperparameters_fit': {'num_boost_round': 446},
    'hyperparameters_nondefault': [],
    'ag_args_fit': {'max_memory_usage_ratio': 1.0,
    'max_time_limit_ratio': 1.0,
    'max_time_limit': None,
    'min_time_limit': 0,
    'valid_raw_types': ['bool', 'int', 'float', 'category'],
    'valid_special_types': None,
    'ignored_type_group_special': None,
    'ignored_type_group_raw': None,
    'get_features_kwargs': None,
    'get_features_kwargs_extra': None,
    'predict_1_batch_size': None,
    'temperature_scalar': None},
    'num_features': 17,
    'features': ['KNeighborsUnif_BAG_L1',
    'KNeighborsDist_BAG_L1',

```

```

'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddcf55780>,
'memory_size': 1353353,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True},
'S1F6': {'name': 'S1F6',
'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 5.495210886001587,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.3667716979980469,
'val_score': -430.3510153400956,
'hyperparameters': {'learning_rate': 0.05},
'hyperparameters_fit': {'num_boost_round': 399},
'hyperparameters_nondefault': [],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,

```

```

    'temperature_scalar': None},
    'num_features': 17,
    'features': ['KNeighborsUnif_BAG_L1',
                  'KNeighborsDist_BAG_L1',
                  'LightGBMXT_BAG_L1',
                  'latitude',
                  'longitude',
                  'city',
                  'province',
                  'lease_term',
                  'type',
                  'beds',
                  'baths',
                  'sq_feet',
                  'furnishing',
                  'availability_date',
                  'smoking',
                  'cats',
                  'dogs'],
    'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aee2435a7a0>,
    'memory_size': 1230921,
    'compile_time': None,
    'is_initialized': True,
    'is_fit': True,
    'is_valid': True,
    'can_infer': True},
    'S1F7': {'name': 'S1F7',
              'model_type': 'LGBModel',
              'problem_type': 'regression',
              'eval_metric': 'root_mean_squared_error',
              'stopping_metric': 'root_mean_squared_error',
              'fit_time': 2.807051181793213,
              'num_classes': None,
              'quantile_levels': None,
              'predict_time': 0.02846217155456543,
              'val_score': -400.9667860832463,
              'hyperparameters': {'learning_rate': 0.05},
              'hyperparameters_fit': {'num_boost_round': 67},
              'hyperparameters_nondefault': [],
              'ag_args_fit': {'max_memory_usage_ratio': 1.0,
                              'max_time_limit_ratio': 1.0,
                              'max_time_limit': None,
                              'min_time_limit': 0,
                              'valid_raw_types': ['bool', 'int', 'float', 'category'],
                              'valid_special_types': None,
                              'ignored_type_group_special': None,

```



```

'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},
'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddcf55960>,
'memory_size': 218244,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True},
'S1F8': {'name': 'S1F8',
'model_type': 'LGBModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 2.716935157775879,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.017421245574951172,
'val_score': -345.3464214922553,
'hyperparameters': {'learning_rate': 0.05},
'hyperparameters_fit': {'num_boost_round': 57},
'hyperparameters_nondefault': [],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,

```

```

'min_time_limit': 0,
'valid_raw_types': ['bool', 'int', 'float', 'category'],
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None},
'num_features': 17,
'features': ['KNeighborsUnif_BAG_L1',
'KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'latitude',
'longitude',
'city',
'province',
'lease_term',
'type',
'beds',
'baths',
'sq_feet',
'furnishing',
'availability_date',
'smoking',
'cats',
'dogs'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aeddcf557b0>,
'memory_size': 185967,
'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True}}},
'WeightedEnsemble_L3': {'name': 'WeightedEnsemble_L3',
'model_type': 'WeightedEnsembleModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 0.048243045806884766,
'num_classes': None,
'quantile_levels': None,
'predict_time': 0.0012357234954833984,
'val_score': -414.8760953952414,
'hyperparameters': {'use_orig_features': False,
'max_base_models': 25,

```

```

    'max_base_models_per_type': 5,
    'save_bag_folds': True},
    'hyperparameters_fit': {},
    'hyperparameters_nondefault': ['save_bag_folds'],
    'ag_args_fit': {'max_memory_usage_ratio': 1.0,
    'max_time_limit_ratio': 1.0,
    'max_time_limit': None,
    'min_time_limit': 0,
    'valid_raw_types': None,
    'valid_special_types': None,
    'ignored_type_group_special': None,
    'ignored_type_group_raw': None,
    'get_features_kwargs': None,
    'get_features_kwargs_extra': None,
    'predict_1_batch_size': None,
    'temperature_scalar': None,
    'drop_unique': False},
    'num_features': 4,
    'features': ['LightGBMXT_BAG_L1',
    'LightGBM_BAG_L2',
    'KNeighborsDist_BAG_L1',
    'LightGBMXT_BAG_L2'],
    'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aedd7526e0>,
    'memory_size': 3143,
    'compile_time': None,
    'is_initialized': True,
    'is_fit': True,
    'is_valid': True,
    'can_infer': True,
    'bagged_info': {'child_model_type': 'GreedyWeightedEnsembleModel',
    'num_child_models': 1,
    'child_model_names': ['S1F1'],
    '_n_repeats': 1,
    '_k_per_n_repeat': [1],
    '_random_state': 3,
    'low_memory': False,
    'bagged_mode': False,
    'max_memory_size': 3143,
    'min_memory_size': 3143,
    'child_hyperparameters': {'ensemble_size': 25, 'subsample_size': 1000000},
    'child_hyperparameters_fit': {'ensemble_size': 10},
    'child_ag_args_fit': {'max_memory_usage_ratio': 1.0,
    'max_time_limit_ratio': 1.0,
    'max_time_limit': None,
    'min_time_limit': 0,
    'valid_raw_types': None,

```

```

'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None,
'drop_unique': False}},
'stacker_info': {'num_base_models': 4,
'base_model_names': ['KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'LightGBMXT_BAG_L2',
'LightGBM_BAG_L2']},
'children_info': {'S1F1': {'name': 'S1F1',
'model_type': 'GreedyWeightedEnsembleModel',
'problem_type': 'regression',
'eval_metric': 'root_mean_squared_error',
'stopping_metric': 'root_mean_squared_error',
'fit_time': 0.048243045806884766,
'num_classes': None,
'quantile_levels': None,
'predict_time': None,
'val_score': None,
'hyperparameters': {'ensemble_size': 25, 'subsample_size': 1000000},
'hyperparameters_fit': {'ensemble_size': 10},
'hyperparameters_nondefault': [],
'ag_args_fit': {'max_memory_usage_ratio': 1.0,
'max_time_limit_ratio': 1.0,
'max_time_limit': None,
'min_time_limit': 0,
'valid_raw_types': None,
'valid_special_types': None,
'ignored_type_group_special': None,
'ignored_type_group_raw': None,
'get_features_kwargs': None,
'get_features_kwargs_extra': None,
'predict_1_batch_size': None,
'temperature_scalar': None,
'drop_unique': False},
'num_features': 4,
'features': ['KNeighborsDist_BAG_L1',
'LightGBMXT_BAG_L1',
'LightGBMXT_BAG_L2',
'LightGBM_BAG_L2'],
'feature_metadata':
<autogluon.common.features.feature_metadata.FeatureMetadata at 0x7aedd752740>,
'memory_size': 5477,

```

```

'compile_time': None,
'is_initialized': True,
'is_fit': True,
'is_valid': True,
'can_infer': True,
'model_weights': {'KNeighborsDist_BAG_L1': 0.1,
'LightGBMXT_BAG_L1': 0.4,
'LightGBMXT_BAG_L2': 0.3,
'LightGBM_BAG_L2': 0.2}}}},
'model_info_failures': {}

```

```

[17]: training_summary = predictor_price.fit_summary()
training_summary

```

\*\*\* Summary of fit() \*\*\*

Estimated performance of each model:

	model	score_val	eval_metric	pred_time_val
fit_time	pred_time_val_marginal	fit_time_marginal	stack_level	can_infer
fit_order				
0	WeightedEnsemble_L3	-414.876095	root_mean_squared_error	132.505690
375.993751		0.001236	0.048243	3
7				
1	WeightedEnsemble_L2	-423.163408	root_mean_squared_error	118.358258
282.499167		0.000892	0.023722	2
4				
2	LightGBMXT_BAG_L2	-432.884149	root_mean_squared_error	131.426344
345.975757		12.976854	63.470796	2
5				
3	LightGBM_BAG_L2	-438.084324	root_mean_squared_error	119.527600
312.474711		1.078110	29.969751	2
6				
4	LightGBMXT_BAG_L1	-442.089932	root_mean_squared_error	118.265830
282.429904		118.265830	282.429904	1
3				
5	KNeighborsDist_BAG_L1	-652.271480	root_mean_squared_error	0.091537
0.045541		0.091537	0.045541	1
2				
6	KNeighborsUnif_BAG_L1	-774.186518	root_mean_squared_error	0.092124
0.029516		0.092124	0.029516	1
1				

Number of models trained: 7

Types of models trained:

```

{'StackerEnsembleModel_KNN', 'StackerEnsembleModel_LGB',
'WeightedEnsembleModel'}

```

Bagging used: True (with 8 folds)

Multi-layer stack-ensembling used: True (with 3 levels)

Feature Metadata (Processed):

```
(raw dtype, special dtypes):
('category', []) : 12 | ['city', 'province', 'lease_term', 'type', 'beds', ...]
('float', [])    : 2 | ['latitude', 'longitude']
Plot summary of models saved to file: agModels-predictpriceSummaryOfModels.html
*** End of fit() summary ***
```

```
[17]: {'model_types': {'KNeighborsUnif_BAG_L1': 'StackerEnsembleModel_KNN',
  'KNeighborsDist_BAG_L1': 'StackerEnsembleModel_KNN',
  'LightGBMXT_BAG_L1': 'StackerEnsembleModel_LGB',
  'WeightedEnsemble_L2': 'WeightedEnsembleModel',
  'LightGBMXT_BAG_L2': 'StackerEnsembleModel_LGB',
  'LightGBM_BAG_L2': 'StackerEnsembleModel_LGB',
  'WeightedEnsemble_L3': 'WeightedEnsembleModel'},
'model_performance': {'KNeighborsUnif_BAG_L1': -774.186517755969,
  'KNeighborsDist_BAG_L1': -652.2714801027624,
  'LightGBMXT_BAG_L1': -442.0899321234507,
  'WeightedEnsemble_L2': -423.163408198703,
  'LightGBMXT_BAG_L2': -432.88414922328326,
  'LightGBM_BAG_L2': -438.08432384123813,
  'WeightedEnsemble_L3': -414.8760953952414},
'model_best': 'WeightedEnsemble_L3',
'model_paths': {'KNeighborsUnif_BAG_L1': ['KNeighborsUnif_BAG_L1'],
  'KNeighborsDist_BAG_L1': ['KNeighborsDist_BAG_L1'],
  'LightGBMXT_BAG_L1': ['LightGBMXT_BAG_L1'],
  'WeightedEnsemble_L2': ['WeightedEnsemble_L2'],
  'LightGBMXT_BAG_L2': ['LightGBMXT_BAG_L2'],
  'LightGBM_BAG_L2': ['LightGBM_BAG_L2'],
  'WeightedEnsemble_L3': ['WeightedEnsemble_L3']},
'model_fit_times': {'KNeighborsUnif_BAG_L1': 0.029515743255615234,
  'KNeighborsDist_BAG_L1': 0.04554128646850586,
  'LightGBMXT_BAG_L1': 282.42990374565125,
  'WeightedEnsemble_L2': 0.023722410202026367,
  'LightGBMXT_BAG_L2': 63.47079634666443,
  'LightGBM_BAG_L2': 29.96975064277649,
  'WeightedEnsemble_L3': 0.048243045806884766},
'model_pred_times': {'KNeighborsUnif_BAG_L1': 0.09212350845336914,
  'KNeighborsDist_BAG_L1': 0.0915365219116211,
  'LightGBMXT_BAG_L1': 118.26582980155945,
  'WeightedEnsemble_L2': 0.0008921623229980469,
  'LightGBMXT_BAG_L2': 12.97685432434082,
  'LightGBM_BAG_L2': 1.0781099796295166,
  'WeightedEnsemble_L3': 0.0012357234954833984},
'num_bag_folds': 8,
'max_stack_level': 3,
'model_hyperparams': {'KNeighborsUnif_BAG_L1': {'use_orig_features': True,
  'max_base_models': 25,
  'max_base_models_per_type': 5,
```

```

'save_bag_folds': True,
'use_child_oof': True},
'KNeighborsDist_BAG_L1': {'use_orig_features': True,
    'max_base_models': 25,
    'max_base_models_per_type': 5,
    'save_bag_folds': True,
    'use_child_oof': True},
'LightGBMXT_BAG_L1': {'use_orig_features': True,
    'max_base_models': 25,
    'max_base_models_per_type': 5,
    'save_bag_folds': True},
'WeightedEnsemble_L2': {'use_orig_features': False,
    'max_base_models': 25,
    'max_base_models_per_type': 5,
    'save_bag_folds': True},
'LightGBMXT_BAG_L2': {'use_orig_features': True,
    'max_base_models': 25,
    'max_base_models_per_type': 5,
    'save_bag_folds': True},
'LightGBM_BAG_L2': {'use_orig_features': True,
    'max_base_models': 25,
    'max_base_models_per_type': 5,
    'save_bag_folds': True},
'WeightedEnsemble_L3': {'use_orig_features': False,
    'max_base_models': 25,
    'max_base_models_per_type': 5,
    'save_bag_folds': True}},
'leaderboard':
    model    score_val    eval_metric
pred_time_val \
0    WeightedEnsemble_L3 -414.876095    root_mean_squared_error    132.505690
1    WeightedEnsemble_L2 -423.163408    root_mean_squared_error    118.358258
2    LightGBMXT_BAG_L2 -432.884149    root_mean_squared_error    131.426344
3    LightGBM_BAG_L2 -438.084324    root_mean_squared_error    119.527600
4    LightGBMXT_BAG_L1 -442.089932    root_mean_squared_error    118.265830
5    KNeighborsDist_BAG_L1 -652.271480    root_mean_squared_error    0.091537
6    KNeighborsUnif_BAG_L1 -774.186518    root_mean_squared_error    0.092124

    fit_time    pred_time_val_marginal    fit_time_marginal    stack_level \
0    375.993751    0.001236    0.048243    3
1    282.499167    0.000892    0.023722    2
2    345.975757    12.976854    63.470796    2
3    312.474711    1.078110    29.969751    2
4    282.429904    118.265830    282.429904    1
5    0.045541    0.091537    0.045541    1
6    0.029516    0.092124    0.029516    1

can_infer    fit_order

```

```

0      True      7
1      True      4
2      True      5
3      True      6
4      True      3
5      True      2
6      True      1 }

```

```

[18]: # Calculate feature importance
feature_importance = predictor_price.feature_importance(data=test_data,
↳time_limit=600)
print(feature_importance)

```

	importance	stddev	p_value	n	p99_high	p99_low
latitude	197.625936	NaN	NaN	1	NaN	NaN
longitude	180.109537	NaN	NaN	1	NaN	NaN
baths	167.315788	NaN	NaN	1	NaN	NaN
beds	156.030511	NaN	NaN	1	NaN	NaN
type	152.126349	NaN	NaN	1	NaN	NaN
furnishing	76.471019	NaN	NaN	1	NaN	NaN
lease_term	45.004177	NaN	NaN	1	NaN	NaN
city	44.273299	NaN	NaN	1	NaN	NaN
availability_date	37.439987	NaN	NaN	1	NaN	NaN
sq_feet	26.604972	NaN	NaN	1	NaN	NaN
province	23.759757	NaN	NaN	1	NaN	NaN
dogs	22.048391	NaN	NaN	1	NaN	NaN
cats	14.550022	NaN	NaN	1	NaN	NaN
smoking	8.246433	NaN	NaN	1	NaN	NaN

- Positive scores indicate features that help the model predict accurately.
- Features with near-zero or negative scores may be redundant or harmful to the model's performance.

```

[20]: import matplotlib.pyplot as plt

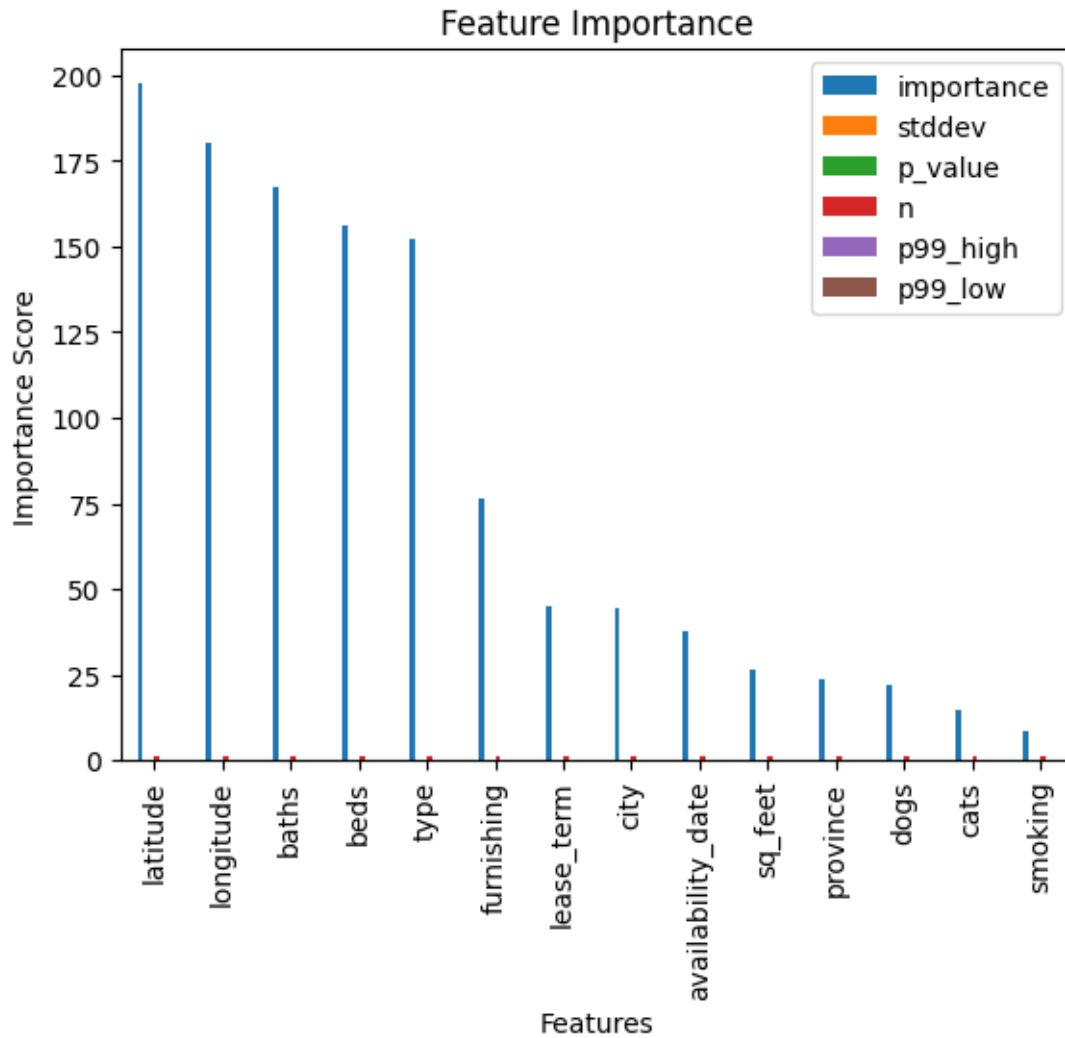
# Sort importance values
sorted_importance = feature_importance.
↳sort_values(by="importance",ascending=False)

# Plot
plt.figure(figsize=(10, 6))
sorted_importance.plot(kind='bar')
plt.title('Feature Importance')
plt.xlabel('Features')
plt.ylabel('Importance Score')
plt.show()

```

<Figure size 1000x600 with 0 Axes>






```
[23]: # Residual Plot
y_pred = predictor_price.predict(test_data.drop(columns=[target]))
y_true = test_data[target]

# Residuals
residuals = y_true - y_pred

plt.figure(figsize=(10, 6))
plt.scatter(y_pred, residuals, alpha=0.5)
plt.axhline(0, color='red', linestyle='--')
plt.title('Residual Plot')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.show()
```



```
[ ]:
```

```
[ ]: # from autogluon.interpret import TabularExplainer


# explainer = TabularExplainer(predictor_price, test_data)
# shap_values = explainer.shap_values()
# explainer.plot_shap_summary(shap_values)
```

```
[24]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

mae = mean_absolute_error(y_true, y_pred)
rmse = mean_squared_error(y_true, y_pred, squared=False)
r2 = r2_score(y_true, y_pred)

metrics = {'MAE': mae, 'RMSE': rmse, 'R2': r2}
plt.bar(metrics.keys(), metrics.values())
plt.title('Model Performance Metrics')
plt.ylabel('Metric Value')
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:483:
FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in
1.6. To calculate the root mean squared error, use the
function 'root_mean_squared_error'.
  warnings.warn(
```



```
[2]: # !pip install shap
```

```
Requirement already satisfied: shap in /usr/local/lib/python3.10/dist-packages (0.46.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from shap) (1.26.4)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from shap) (1.12.0)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from shap) (1.4.0)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from shap) (2.2.2)
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.10/dist-packages (from shap) (4.66.6)
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.10/dist-packages (from shap) (24.2)
Requirement already satisfied: slicer==0.0.8 in /usr/local/lib/python3.10/dist-packages (from shap) (0.0.8)
Requirement already satisfied: numba in /usr/local/lib/python3.10/dist-packages (from shap) (0.60.0)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dist-
```

```
packages (from shap) (3.1.0)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in
/usr/local/lib/python3.10/dist-packages (from numba->shap) (0.43.0)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas->shap) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
packages (from pandas->shap) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-
packages (from pandas->shap) (2024.2)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-
packages (from scikit-learn->shap) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (3.5.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-
packages (from python-dateutil>=2.8.2->pandas->shap) (1.16.0)
```

[ ]: